

Innovation-Normalized Detection in Compound-Gaussian Clutter: Exact Laws, Conformal Thresholds and Certified Integration under Pulsed Interference

Muhammad Umair Raza, Sohail Ahmed, and Ammad Ahmed

Abstract—Constant-false-alarm-rate (CFAR) laws assume independent exponential samples, which correlated, textured sea clutter does not provide. Whitening each range cell’s slow-time history with an autoregressive model and normalizing the tested innovation by the cell’s own innovation power restores these laws under the true compound-Gaussian model, at a whitened signal-to-clutter ratio. We derive a post-onset law: a persistent target masks itself after a predictable number of looks, and a slow-time guard of Δ pulses delays this by Δ looks. A conformal threshold on a worst-case bound certifies clipped and binary integration against pulsed interference of any power, if interfered pulses are clutter-independent and dominated by the calibration hit model. On IPIX, conformal thresholds keep 12 of 13 statistics within 0.73–1.47 times the design rate down to 10^{-4} , whereas model-based laws of 10 of 11 whitened and ANMF statistics exceed twice that rate, up to 60-fold. At a design rate of 10^{-2} , an eight-pulse guard raises the mean per-look detection probability of an injected persistent 10 dB target from 0.24 to 0.82, except at the clutter Doppler. At that design rate, with injected interference on 5% of pulses, certified integration stays below design and detects that target with probability 0.64, whereas uncertified non-coherent integration exceeds design up to 23-fold. On NetRAD S-band clutter, unused in development, the guard and the injected-interference certificate behave as on IPIX and conformal rates stay within 1.36 times design at 10^{-3} , but 5 of 13 statistics exceed twice design at 10^{-4} . Injected bursts defeat every remedy tested.

Index Terms—Compound-Gaussian clutter, conformal prediction, constant false alarm rate, pulse integration, pulsed interference, sea clutter.

I. INTRODUCTION

CFAR design rests on a few exact laws: the false-alarm and detection probabilities of cell-averaging (CA) and order-statistic (OS) detectors, and of non-coherent, binary and clipped pulse integration, all derived for independent exponential samples in thermal noise [1]–[5]. Sea clutter violates this premise: a slowly varying texture scales locally Gaussian speckle, the speckle is correlated from pulse to pulse according to its Doppler spectrum, and thermal noise adds a floor [6]–[8]. A threshold set by the thermal-noise laws then neither

keeps the false-alarm rate constant across clutter power nor places it correctly in the tail.

A detector that tests each range cell along slow time faces three further problems. First, its threshold must hold far in the tail. Adaptive detectors restore scale invariance with a covariance estimated from secondary data [7], [9]–[11] or with an AR whitening filter [12], [13], but take their thresholds from Gaussian laws or simulation, whose tail accuracy on real clutter is rarely measured. Learned and resampling sea-clutter detectors set thresholds empirically [14]–[18], without finite-sample guarantees or performance laws. Second, a detector that normalizes a cell by its own past can mask a target that persists into that past, as clutter-map detectors do [19], [20]. Third, pulsed interference from other radars, jammers or the chirp crossings of frequency-modulated continuous-wave (FMCW) radars corrupts a few pulses with unknown power, and with them both the test statistic and its reference.

We address these problems with one construction (Fig. 1). Each cell’s slow-time history is whitened by an AR model fitted on clutter training data, and the tested innovation is normalized by the innovation power of the cell’s own history. At the true AR coefficient and without thermal noise, the innovation powers are independent exponentials whose common scale, the texture, cancels in the normalization, so the thermal-noise laws return with the signal-to-clutter ratio (SCR) multiplied by the whitening factor. The per-look score is a one-pulse parametric adaptive matched filter (PAMF) test [12]. We call it IN-ARCP (innovation-normalized autoregressive conformal prediction), because its exceedance region is the complement of a conformal prediction disc for the next sample [21]. The contributions are as follows.

- 1) **Laws** (Section IV). Whitening and self-normalization return the CA, OS, integration and normalized-matched-filter laws with the whitened SCR. A closed-form detection law after onset gives a *visibility horizon*, the number of looks after which a strong persistent target has masked itself; a slow-time guard, the analogue of the guard cells of range CFAR [4], moves the horizon by the guard length. We also give the OS counterpart. Monte Carlo confirms the laws under the model; in its main run, all 65 checks of the closed forms lie within 3 standard errors (Supplementary Material).
- 2) **Calibration** (Section IV-F). The known finite-sample law

M. U. Raza is with the Department of Avionics Engineering, College of Aeronautical Engineering, National University of Sciences and Technology (NUST), Risalpur, Pakistan (e-mail: uraza@cae.nust.edu.pk). Corresponding author: M. U. Raza.

S. Ahmed is with MathWorks, Natick, MA, USA.

A. Ahmed is with the National Disaster Management Authority (NDMA), Islamabad, Pakistan.

of split-conformal thresholds [22] fixes the calibration size. On IPIX these thresholds hold the false-alarm rate to 10^{-4} , and on NetRAD S-band sea clutter to 10^{-3} , where the model-based laws of most whitened detectors and of the adaptive normalized matched filter (ANMF) fail.

- 3) **Certified integration under pulsed interference** (Theorem 1). For clipped [5] and binary integration, a conformal threshold on a worst-case bound over interference amplitudes bounds the false-alarm probability for interference of any power. The pattern of interfered pulses must be independent of the clutter and dominated, in the inclusion order, by the hit model used in calibration.

We test all three on IPIX sea clutter, on NetRAD S-band sea clutter unused in development, and on 77 GHz FMCW data with real interference, against the PAMF, the ANMF, power detectors, certified trimming and pulse blanking. We do not claim a more sensitive detector, since the whitening gain is shared with the PAMF. The contribution is a detection framework whose false-alarm rate is calibrated on real clutter, whose performance is predicted by exact laws, and whose integration can be certified under interference.

II. RELATED WORK

Normalized detection in compound-Gaussian clutter. Scale invariance of normalized statistics underlies texture-CFAR detection [7], [8], [10]. The ANMF, or adaptive coherence estimator (ACE), normalizes by a covariance estimated from secondary data with the sample covariance matrix (SCM) [10] or the fixed-point (Tyler) estimator [11]; the PAMF and normalized PAMF whiten with an AR filter [12], [13]. Detection through the prediction error of a clutter model is classical, including on IPIX [23]. Targets in the reference window mask CA detectors, and they mask OS detectors once they occupy more than $m - k$ reference cells [2], [3]. Clutter-map detectors filter a cell's own past power [19], and their self-masking by persistent targets has been analyzed and countered by censoring [20]. We found no such analysis for the AR-innovation score of a cell's own history.

Data-driven thresholds. Split conformal prediction gives finite-sample marginal coverage for exchangeable data [21], [22], and a conformal p -value is a calibrated test [24]. Against the empirical thresholds of learned and AR-sieve detectors, a conformal threshold adds a finite-sample law, and innovation normalization adds a score that is nearly pivotal where clutter dominates thermal noise.

Integration under pulsed interference. Non-coherent, binary and censored integration in random impulse interference was analyzed in [25], and clipping the per-pulse statistics is a robust remedy [5]; for automotive FMCW interference, fast-time zeroing and reconstruction are standard [26], [27]. Robust conformal prediction certifies coverage under bounded or sparse perturbations [28], [29], randomized ablation certifies against sparse perturbations of unbounded amplitude [30], and calibrating over masks that contain the test mask handles missing values [31]. Theorem 1 applies the mask device to pulsed interference, where order statistics give a closed-form bound over all amplitudes.

III. MODEL AND DETECTORS

A. Compound-Gaussian autoregressive episodes with thermal noise

An *episode* is $X = (H_1, \dots, H_m, Y) \in \mathbb{C}^{m+1}$: m consecutive slow-time samples of one range cell (the history) followed by the next sample Y . Conditional on a latent texture,

$$X | c \sim \mathcal{CN}(0, \nu(cR + I)), \quad (1)$$

where R is the unit-diagonal speckle correlation, $\nu \geq 0$ the thermal-noise power and $c > 0$ the clutter-to-noise ratio (CNR), which carries the texture. Letting $\nu \rightarrow 0$ with νc fixed gives the compound-Gaussian model [7], [8]; the additive term is the thermal noise of K+noise clutter [6], [32]. The speckle is a stationary complex AR(p) process. Episodes from disjoint times or cells are treated as exchangeable sampling units, that is, their joint distribution is assumed unchanged by reordering. A *look* is one test of a pulse by the per-look score; look ℓ is the ℓ -th pulse after a target's onset. A *dwell* is K consecutive tested pulses after the history.

B. Whitened innovations and detectors

Every whitened detector uses the innovations of an AR model fitted by least squares on training episodes. For AR(1) with coefficient r , radially clipped to $|r| \leq 0.98$ (the clip binds in 5 of 56 IPIX units), $\varepsilon_1 = \sqrt{1 - |r|^2}h_1$ and $\varepsilon_j = h_j - rh_{j-1}$; for AR(p), the innovations are the $m - p$ conditional ones.

- *Per-look score (IN-ARCP).* $|Y - rh_m|/s_r(h)$ with $s_r(h)^2 = m^{-1} \sum_j |\varepsilon_j|^2$. Its exceedance of the threshold $\hat{q}$ (Section III-C) is the complement of the prediction disc $\{y : |y - rh_m| \leq \hat{q}s_r(h)\}$. With a one-pulse steering vector, the PAMF statistic normalized by the history innovations equals, by construction, the squared AR(p) score (IN-AR(p)). With a *guard* of Δ pulses, the scale is computed from the m samples that end Δ pulses before h_m , while the center still uses h_m (Fig. 1); $\Delta = 0$ is IN-ARCP.
- *OS score.* The scale is the k -th smallest of the m history innovation powers.
- *Dwell detectors.* Over K dwell pulses after an m -pulse history, non-coherent integration sums the dwell innovation powers against their history mean. Clipped integration computes $\sum_{i=1}^K \min\{|\varepsilon_{m+i}|^2/s_{k_H}^2, \kappa\}$ with clip level κ , and binary integration counts the dwell powers $|\varepsilon_{m+i}|^2/s_{k_H}^2$ above η and thresholds the count; $s_{k_H}^2$ is the k_H -th smallest history innovation power. PAMF-H is the coherent PAMF over the dwell, normalized by the history innovations. Pulse blanking, a non-certified alternative, discards an innovation power, and the innovation after it, when the power exceeds both ten times the history median and four times a robust level of its block. That level is the history median for history terms and the third-smallest dwell power for dwell terms. The remaining dwell powers are integrated.
- *Self-normalized whitened Doppler (parametric ANMF, P-ANMF).* $\max_j |\text{DFT}_j(\varepsilon)|^2/(N\|\varepsilon\|^2)$ over N dwell innovations: the ANMF with an AR covariance, normalized by the dwell itself rather than by the history.




Fig. 1. The detector in the receive chain. It takes the place of the range CFAR stage after range processing. Each range bin is tested along slow time: the tested innovation is normalized by the innovation power of a scale window of m pulses that ends Δ pulses before the tested pulse Y , and compared with a threshold calibrated on clean clutter. Neighboring bins are not used.

- **Noise-aware (NA) score.** The Gaussian predictive mean and variance of (1) given h , with (R, ν) and a nonparametric texture law fitted by marginal likelihood and a per-episode CNR [33]. It reduces to IN-ARCP as $\nu \rightarrow 0$ (Supplement).

C. Conformal thresholds

For a score T and a design false-alarm rate α , the threshold $\hat{q}$ is the k_α -th smallest value of T over n clean calibration episodes, with $k_\alpha = \lceil (n+1)(1-\alpha) \rceil$. For exchangeable calibration and test episodes, the false-alarm probability is then at most α [21], [24]. Section IV-F gives its distribution given the calibration set.

IV. LAWS OF WHITENED-INNOVATION DETECTION

The laws follow a target through the detector, from the tested pulse at onset to a target that persists into the history and a dwell of several pulses; the calibration and certification of the threshold close the section.

A. The exact law at onset

Proposition 1. Let $X \sim \mathcal{CN}(0, \Sigma)$, $\Sigma \succ 0$, fix a linear center $a^\top H$ and a scale $s(H)^2 = H^\top Q H / m$ with $Q \succeq 0$, $Q \neq 0$, and let λ_k be the nonzero eigenvalues, with multiplicity, of $\Sigma^{1/2} M_q \Sigma^{1/2}$, $M_q = \bar{b} \bar{b}^\top - \frac{q}{m} \text{diag}(Q, 0)$, $\bar{b} = e_{m+1} - (a^\top, 0)^\top$. Exactly one eigenvalue λ_+ is positive, and

$$G_\Sigma(q) := \mathbb{P}\{|Y - a^\top H| \leq q s(H)\} = 1 - \prod_{k \neq +} \frac{\lambda_+}{\lambda_+ - \lambda_k}. \quad (2)$$

The proof (Appendix A) holds for any multiplicities. Equation (2) specializes classical quadratic-form theory [34]. At the true AR(1) coefficient it is the CA law $1 - (1 + q^2/m)^{-m}$ [1], [3] applied to innovations. It evaluates coverage, the probability that the next sample falls inside the prediction disc, at any linear center and quadratic scale, texture and noise level, without simulation.

Corollary 1. With $\nu = 0$, (2) does not depend on the texture. With $\nu > 0$ texture invariance is lost [32], [35]; the CNR-conditional coverage is $c \mapsto G_{cR+I}(\hat{q})$, which tends to the pure-clutter value as $c \rightarrow \infty$ and to $G_I(\hat{q})$ as $c \rightarrow 0$.

Corollary 2. A Swerling-1 target $\sqrt{SP} g v$ (P the local clutter-plus-noise power, S the SCR relative to it, $g \sim \mathcal{CN}(0, 1)$, v the target signature over the episode) is detected

with $P_d = 1 - G_{\Sigma+SPvv^\top}(\hat{q})$. For a target in Y only, pure AR(1) clutter and the true coefficient,

$$P_d = \left(1 + \frac{q^2}{m} \frac{1 - |\rho|^2}{1 - |\rho|^2 + S}\right)^{-m}, \quad (3)$$

the CA-CFAR law with the SCR multiplied by the whitening factor $1/(1 - |\rho|^2)$.

Because the innovation powers are then independent exponentials, the OS-CFAR law also holds at onset.

Corollary 3. With $\nu = 0$, AR(1) speckle and $r = \rho$, the OS score with the k -th smallest of m innovation powers has $\mathbb{P}\{\text{exceedance}\} = \prod_{i=0}^{k-1} (m-i)/(m-i+q^2)$ for every texture, the OS-CFAR law [2]; a target in Y only is detected with this probability, q^2 replaced by $q^2/\{1 + S/(1 - |\rho|^2)\}$.

B. Thermal noise caps the whitening gain

Proposition 2. Under (1) with AR(1) speckle, testing the prediction error against its own variance gains, over power-only detection and for a long history, $G_{\text{IN}}(c) = (1+c)/\{c(1+|r|^2 - 2\Re \bar{r} \rho) + 1 + |r|^2\}$ for the IN-ARCP center and $G_{\text{opt}}(c) = (1+c)/V(c) \geq \max\{1, G_{\text{IN}}\}$ for the optimal predictor; with $V = \frac{1}{2}(\xi + \sqrt{\xi^2 - 4|\rho|^2})$ and $\xi = 1 + |\rho|^2 + c(1 - |\rho|^2)$.

The prediction error $Y - rH_m$ is a weighted two-pulse canceller, and G_{IN} is its clutter attenuation, which equals its improvement factor for a target in Y alone [4]. In noise-dominated cells, first-order normalization loses up to $10 \log_{10}(1 + |\rho|^2)$ dB, which the NA score avoids (proof in the Supplementary Material).

C. Persistent targets

A target that remains in the beam after its onset enters the history that normalizes it; in the limit of a strong target present in every sample, the score cannot detect it.

Proposition 3. For a persistent target $A(e^{j\omega t})_{t=1}^{m+1}$ in every sample, the IN-ARCP score tends as $|A| \rightarrow \infty$ to $T_\infty = \sqrt{m}|1 - re^{-j\omega}|/\{(1 - |r|^2) + (m-1)|1 - re^{-j\omega}|^2\}^{1/2} \leq \sqrt{m/(m-1)}$, so $P_d \rightarrow 0$ whenever $\hat{q} > \sqrt{m/(m-1)}$.

To follow a target that starts at a known pulse, we fix the model for the rest of this section: the speckle is AR(1) with coefficient ρ , $\nu = 0$ and $r = \rho$ (Theorem 1 needs none of these assumptions). The innovations $\varepsilon_1 = \sqrt{1 - |\rho|^2} x_1$ and $\varepsilon_t = x_t - \rho x_{t-1}$ are then independent $\mathcal{CN}(0, \sigma_e^2)$ given the texture, and every statistic below is homogeneous of degree zero, so the texture cancels. A Swerling-1 target $\sqrt{SP} g e^{j\omega t}$, $g \sim \mathcal{CN}(0, 1)$ common to all pulses, that starts at pulse t_0 contributes $g\sqrt{S_w}$ to ε_{t_0} and $g\sqrt{S_w} d(\omega) e^{j\omega t}$ to ε_t , $t > t_0$, where

$$S_w = \frac{S}{1 - |\rho|^2}, \quad |d(\omega)|^2 = |1 - \rho e^{-j\omega}|^2, \quad (4)$$

$$S_w |d(\omega)|^2 = S w(\omega), \quad w(\omega) = \frac{P}{S_c(\omega)},$$

and $S_c(\omega) = P(1 - |\rho|^2)/|1 - \rho e^{-j\omega}|^2$ is the clutter spectrum. Whitening thus yields independent exponential innovation




Fig. 2. Closed forms (lines) against Monte Carlo (markers; 10^5 episodes per point, 5×10^4 in (b)) at the true coefficient, AR(1) compound-Gaussian clutter with $\rho = 0.93e^{-0.2j}$, lognormal texture, $m = 16$, $P_{fa} = 0.01$. (a) Proposition 4; (b) horizon, Corollary 4; (c) OS scales (quadrature); (d) Theorem 1, P_{fa} against the hit rate for Gaussian interference 30 dB above the local power, Monte Carlo only (open triangles: no event in 10^5).

powers. It multiplies the SCR by $1/(1 - |\rho|^2)$ in the onset pulse and by $w(\omega)$ after it; for AR(p) speckle, $d(\omega)$ is the whitening filter's response. Self-masking is classical for clutter maps [19], [20]; to our knowledge, the post-onset law with its horizon, the OS law under persistence and the certificate are new in this form.

D. After the onset: a rank-one law and a visibility horizon

Proposition 4. Let the innovation-normalized score with m history innovations be thresholded at q , and let a Swerling-1 target have whitened energy $\mathcal{E}_0$ in the tested innovation and $\mathcal{E}_H$ in the history innovations. With $\beta = q^2/m$, let $\lambda_{\pm}$ be the roots of

$$\lambda^2 - (1 + \mathcal{E}_0 - \beta - \beta\mathcal{E}_H)\lambda - \beta(1 + \mathcal{E}_0 + \mathcal{E}_H) = 0.$$

Then $\lambda_+ > 0 > \lambda_-$ and

$$P_d = \frac{\lambda_+}{\lambda_+ - \lambda_-} \left(1 + \frac{\beta}{\lambda_+}\right)^{-(m-1)}. \quad (5)$$

For a persistent target observed at look $1 \leq \ell \leq m-1$, $\mathcal{E}_0 = S_w |d(\omega)|^2$ and $\mathcal{E}_H = S_w \{1 + (\ell-1)|d(\omega)|^2\}$; at onset ($\ell = 0$), $\mathcal{E}_0 = S_w$, $\mathcal{E}_H = 0$, and (5) reduces to (3).

Corollary 4. As $S \rightarrow \infty$, the look- ℓ detection probability tends to one if $\ell < \ell^*$ and to zero if $\ell > \ell^*$, where

$$\ell^*(\omega) = 1 + \frac{m}{q^2} - \frac{1}{|d(\omega)|^2}. \quad (6)$$

Since $\ell^* < 1 + m/q^2$, a strong target is kept for fewer than $1 + m/q^2$ looks, and not at all near the clutter Doppler, where $|d(\omega)|^2 \approx (1 - |\rho|^2)$. At $m = 16$ and $P_{fa} = 0.01$, $m/q^2 = 3.0$; for $|\rho| = 0.93$ the horizon is 3.7 looks at the opposite Doppler and negative at the clutter Doppler (Fig. 2a,b). A guard delays the loss. With a guard of Δ pulses, the scale excludes the first Δ target innovations. Equation (5) then holds with ℓ replaced by $\ell - \Delta$ ($\mathcal{E}_H = 0$ for $\ell \leq \Delta$), so the horizon becomes $\Delta + \ell^*$, while the false-alarm law is unchanged.

With the OS scale, the target occupies ℓ history innovations after the onset, and P_d is a two-dimensional integral (quadrature). For a strong target and $\ell \leq m - k$, the contaminated

powers lie above the scale, which becomes the k -th smallest of $m - \ell$ clean powers, and

$$P_d \simeq \prod_{i=0}^{k-1} \frac{m - \ell - i}{m - \ell - i + q^2/(1 + S_w |d|^2)}, \quad \ell \leq m - k, \quad (7)$$

while $P_d \rightarrow 0$ for $\ell > m - k$ (Fig. 2c).

E. Dwell integration and self-normalized Doppler

Over a dwell, non-coherent integration of K innovation powers against the mean of m history powers, at threshold θ and with $\tau = \theta/m$, has

$$P_{fa} = \sum_{i=0}^{K-1} \binom{m+i-1}{i} \frac{\tau^i}{(1+\tau)^{m+i}}, \quad (8)$$

and a persistent Swerling-1 target of whitened dwell energy $\mathcal{E}_D = S_w \{1 + (K-1)|d|^2\}$, with $\delta = 1 + \mathcal{E}_D$ and $\tau' = \tau(1 - 1/\delta)$, has

$$P_d = \sum_{i=0}^{K-2} \binom{m+i-1}{i} \frac{\tau^i}{(1+\tau)^{m+i}} + \left(1 - \frac{1}{\delta}\right)^{-(K-1)} \times \left[\left(1 + \frac{\tau}{\delta}\right)^{-m} - \sum_{i=0}^{K-2} \binom{m+i-1}{i} \frac{\tau'^i}{(1+\tau)^{m+i}} \right], \quad (9)$$

the CA non-coherent-integration law [3], [4] with the whitened energy. Power integration of K correlated pulses gains only $K_{\text{eff}} = K^2 / \sum_{i,j} |\rho|^{2|i-j|}$ independent samples (1.41 for $K = 8$ at $|\rho| = 0.93$).

The P-ANMF divides the periodogram of the N whitened innovations by their energy. For a known on-grid Doppler, the statistic is Beta(1, $N-1$) under H_0 for every texture, and

$$P_d = \left[1 + \frac{P_{fa}^{-1/(N-1)} - 1}{1 + NS_w(\omega)} \right]^{-(N-1)}, \quad (10)$$

the normalized-matched-filter law [7], [10] with the whitened SCR. Maximizing over the N DFT bins gives Fisher's g statistic [36]; its null law and, by inclusion-exclusion over the bins with one non-central ordinate, its detection law are elementary (Supplementary Material). Because the dwell normalizes itself, a target that persists through the history is not masked, whereas Proposition 3 rules this out for history normalization.

F. Conformal calibration

For i.i.d. continuous scores, the false-alarm probability given the n calibration scores is $\text{Beta}(n+1-k_\alpha, k_\alpha)$ [22]. With ties, as for clipped or binary statistics, $\mathbb{P}\{T > \hat{q}\} \leq \alpha$ still holds, conservatively. Because k_α is a ceiling, $\mathbb{P}\{P_{\text{fa}} > 2\alpha\}$ is not monotone in n . It first falls below 0.05 at $n = 149$, 1,497 and 14,978 for $\alpha = 10^{-2}, 10^{-3}, 10^{-4}$, and it stays below for every n beyond about $3.15/\alpha$. Our 10^{-4} calibration sets of 12,100–23,100 episodes give 0.02–0.09. Dependence among episodes reduces the effective n , so these sizes are lower bounds. For a known-scale Swerling-1 detector, the expected P_d under calibration is a ratio of Beta functions (Supplementary Material).

G. Certified integration under pulsed interference

Pulsed interference (asynchronous pulses of another radar, jamming pulses or FMCW chirp crossings) corrupts a few pulses with interference of arbitrary, unknown power. We call such a corrupted pulse a *hit*, to distinguish it from an exceedance, a statistic above its threshold. For the whitened statistics, a hit at pulse t_h corrupts the innovations $t_h, \dots, t_h + p$. The clipped integrator $T = \sum_{i=1}^K \min\{|\varepsilon_{m+i}|^2/s_{k_H}^2, \kappa\}$ and its binary counterpart admit a certificate that does not depend on the interference power.

Theorem 1. *For a clean episode x and a set $\mathcal{J}$ of corrupted innovations, let $W(x, \mathcal{J})$ be T computed with the corrupted history powers set to zero and the corrupted dwell terms set to κ (binary: counted as exceedances). Then $T \leq W(x, \mathcal{J})$ for every value of the interference, and W is nondecreasing in $\mathcal{J}$. Let $\hat{q}$ be the split-conformal quantile of $W(X_i, \mathcal{J}_i)$ over clean calibration episodes X_i , with $\mathcal{J}_i$ drawn from a hit model $\mathcal{D}$ independently of the calibration data. If the clean calibration and test episodes are exchangeable and the test hit set is independent of the clutter and dominated by $\mathcal{D}$ in the inclusion order (there is a coupling $\mathcal{J} \subseteq \tilde{\mathcal{J}} \sim \mathcal{D}$), the false-alarm probability is at most α for any interference amplitudes, including adversarial ones.*

A bound on the hit rate alone does not give this dominance: periodic or bursty interference with the same rate can put more hits into one segment than independent hits do. A trimmed sum, which discards the largest dwell terms, is certified at the same level α when the corrupted dwell terms are set to $+\infty$ in W . Its threshold is finite only if the probability that more dwell terms are corrupted than discarded is below α . We discard the fewest terms for which this probability is at most $\alpha/5$: 4 of eight at a 2% hit rate and $\alpha = 10^{-2}$, since a hit corrupts two AR(1) innovations. Fig. 2d checks the theorem in simulation.

V. EXPERIMENTAL DESIGN

IPIX sea clutter. The 14 openly available Dartmouth 1993 stare sessions of the McMaster IPIX radar [37]: X-band, 1 kHz pulse repetition, 14 range cells of 15 m, six dates with wave heights of 0.8–3.8 m [23]. Preprocessing follows the providers' recipe. Target cells and one guard cell on each side are excluded from clutter; Section VI-C uses the primary target

cell. The two like-polarized channels are analyzed separately. Each session is split into thirds, separated by 2,000-pulse gaps, that serve as training, calibration and test data. One of the three assignments of these roles was used while the methods were developed; the other two give the 56 evaluation units (session $\times$ channel $\times$ assignment), which therefore reuse recordings seen during development and are not an untouched confirmation. Day transfer and two series held out until the end are in the Supplementary Material.

77 GHz FMCW data. The open JKU parking-lot dataset [38]: 76–77 GHz chirp sequences (34 μs chirp period, 128 chirps per frame, 100 frames) recorded by two radar stations observing a static scene of parked cars, with a clean reference run and runs with an unsynchronized interferer, dense (scenario A, 45% of chirps hit) and sparse (scenario C, 4.6%). A further scenario-A run adds a walking person. A chirp counts as a hit when more than 10% of its range bins exceed ten times the reference run's median power.

NetRAD S-band sea clutter. Fourteen recordings of the open NetRAD data [39], [40]: 2.4 GHz, 45 MHz chirp, 1 kHz pulse repetition, monostatic node, Cape Peninsula, 9 June 2011, seven HH and seven VV recordings of 130 s, with the providers' clutter range cells (41–101 per recording). No statistic was computed on these data before their protocol was frozen. Detectors and settings are those used on IPIX, except that the ANMF secondary data come from the cells two to four cells away (15–30 vectors). The 28 units are recording $\times$ assignment, and intervals resample recordings. A pulse is flagged when more than 10% of its range bins exceed ten times the bin's median power over the recording. Whether flagged pulses are interference or range-extended sea spikes is not established.

Protocols. For each study, the methods, settings, metrics and expected outcomes were written down and the file's hash recorded in the code repository before any outcome was computed (an internal record, not an external registry); analyses added afterwards are labeled exploratory. Of the 53 outcomes predicted for the four final studies, 13 were not met and 4, all from the pedestrian study (Section VI-B), were left without a verdict; all are listed, with the protocol deviations, in the Supplementary Material.

Detectors and thresholds. Per-look scores use $m = 16$. Dwell detectors use a 16-pulse history and an 8-pulse dwell that starts at the first full-amplitude target pulse (the same idealization for every detector), the median as OS history scale, clip level $\kappa = 6$ (4 to 18 in a clip-level comparison) and binary threshold $\eta = 6$. The ANMF uses the 8-pulse dwell and secondary vectors from clutter cells at least two cells away, at five time positions (20–35 vectors). Power detectors normalize by the same 16 history samples (CA16) or by the cell's power over ± 1024 pulses (CAloc, non-causal). Conformal thresholds are calibrated on clean episodes from the calibration third. For the low-false-alarm analysis these episodes do not overlap (17,200–23,100 per unit; 12,100–16,300 dwell segments). Model-based thresholds are the laws of Section IV at the fitted model, Gaussian plug-ins, the Kraut–Scharf and fixed-point ANMF laws, Gaussian simulation and an AR residual bootstrap of the sieve type [17].

TABLE I

MEASURED FALSE-ALARM RATE / DESIGN α ON IPIX TEST THIRDS (56 UNITS, POOLED). CONFORMAL: SPLIT-CONFORMAL THRESHOLD. MODEL-BASED: THE RULE IN PARENTHESES AT THE FITTED MODEL. BOLD: OUTSIDE $[0.5, 2]$. DAY-CLUSTER INTERVALS ARE IN THE SUPPLEMENTARY MATERIAL; AT 10^{-4} EACH UNIT CONTRIBUTES 1.2–2.2 EXPECTED FALSE ALARMS.

α	Conformal			Model-based rule		
	10^{-2}	10^{-3}	10^{-4}	10^{-2}	10^{-3}	10^{-4}
<i>Per look</i>						
IN-ARCP (CA law)	1.01	1.03	0.73	0.96	1.02	1.92
residual bootstrap				1.63	2.43	4.27
IN-AR(4) (CA law)	1.00	1.02	0.86	1.03	1.19	2.73
OS scale, $k = 8$ (OS law)	1.04	1.02	0.92	2.15	8.34	45.9
NA-AR(4) (Gaussian)	1.00	1.05	1.04	1.56	2.79	7.46
Power, CA (white law)	1.00	1.00	0.75	0.80	0.56	0.49
Power, OS (white law)	1.01	0.98	0.75	2.07	5.56	18.3
<i>Eight-pulse dwell</i>						
Integration (law (8))	1.02	1.05	1.02	2.02	4.35	14.3
PAMF-H, one Doppler (CA law)	0.99	1.00	0.89	1.97	3.66	9.88
P-ANMF, one bin (Beta)	1.10	1.14	1.10	1.78	3.05	5.47
P-ANMF, max (Fisher's g)	1.04	1.04	1.47	4.36	13.1	49.6
ANMF-SCM (Kraut–Scharf)	1.06	1.04	0.85	2.02	3.57	8.07
ANMF-Tyler (FP law)	1.04	1.05	0.82	1.62	2.66	6.01
Clipped, OS scale (Gaussian sim.)	1.05	0.67	0.03	4.84	17.2	59.9

Targets, interference and statistics. Swerling-1 targets are injected at SCRs from -5 to 25 dB relative to the local power: in the tested pulse (onset), persistent from a given pulse with random Doppler, or rising linearly over L pulses. Pulsed interference is injected into IPIX test segments and secondary data: each pulse is hit independently with probability $p_h \in \{0.02, 0.05\}$, with complex Gaussian amplitude 10 or 30 dB above the local power. On the 77 GHz data, targets are injected in the raw analog-to-digital-converter samples as beat tones and pass through the same processing. Intervals are 95% day-cluster bootstrap intervals (six days, 10,000 resamples). With six clusters they are approximate; leave-one-day-out ranges are in the Supplementary Material.

VI. RESULTS

A. False-alarm control across clutter power and down to 10^{-4}

Table I is the central calibration result. Conformal thresholds keep the measured rate of twelve of the thirteen statistics within 0.98–1.14 times design at 10^{-3} and 0.73–1.47 times at 10^{-4} (IN-ARCP: 0.73 [0.54, 0.99]). The exception is the clipped integrator, which is conservative (0.67 and 0.03) because its statistic saturates at $K\kappa$ and ties at the threshold. At 10^{-2} some ratios lie slightly above one, such as 1.10 [1.01, 1.19] for the one-bin P-ANMF. These exceed sampling noise and suggest that the calibration and test thirds of a session are not exactly exchangeable.

Model-based thresholds, by contrast, fail in the tail. At 10^{-4} , 10 of the 11 model-based laws of the whitened statistics and the ANMF exceed twice the design rate, by up to 60-fold (bold entries in Table I). The ANMF laws fail already at 10^{-3} , at 3.57 (SCM) and 2.66 (Tyler) times design. The Kraut–Scharf law fails 2.8–3.4-fold even in simulated compound-Gaussian clutter (Supplementary Material), so the SCM failure need not stem from the clutter tails. A residual bootstrap

of real innovations still exceeds design 4.27-fold at 10^{-4} , consistent with dependence among an episode's innovations, which independent resampling discards. The only law that stays within twofold is the CA law of the IN-ARCP score (1.02 at 10^{-3} ; 1.92 [1.26, 2.61] at 10^{-4} , still above design), plausibly because its scale averages all 16 history innovations.

The false-alarm rate should also stay constant across clutter power. At a design of 0.01, the conformal false-alarm rate of IN-ARCP ranges from 0.0067 to 0.0118 over texture quintiles, a max/min spread of 1.7 (2.7 at 10^{-4} in the low-false-alarm analysis). Without normalization the spread is 87.5; normalization by the raw history RMS gives 9.8. Normalizing a prediction residual by the raw history RMS therefore does not keep the score pivotal (with a null distribution free of the clutter power), whereas innovation normalization keeps it approximately so where clutter dominates noise. In the larger calibration sets of the low-false-alarm analysis, the spread at 10^{-2} is 1.31 for IN-ARCP and 1.06 for the CA16 power ratio. CA16 is thus more pivotal still, but it needs 7.0 dB more SCR.

On the 77 GHz data, thermal noise breaks the texture invariance as Corollary 1 predicts. From the lowest CNR class to the 10–20 dB class, IN-ARCP coverage at a design of 0.90 rises from 0.897 to 0.962, more than the law's 0.932 (mean absolute deviation over all classes 0.010), and the false-alarm rate at 10^{-2} falls further than predicted (post hoc; Table IV).

B. Detection of new returns

Fig. 3a shows targets that appear in the tested pulse. At $P_d = 0.5$, IN-ARCP needs 7.0 dB [4.3, 9.0] less SCR than power detection on the same pulses (CA16) and 9.8 dB less than the long-window power detector (CALoc; 3.5–13.8 dB by day, post hoc). The exact law of Corollary 2 predicts P_d with a mean absolute error of 0.021 per unit and SCR. It attributes the shortfall from the clutter-only whitening factor (11.2 dB) to thermal noise (Proposition 2). The Gaussian threshold of IN-ARCP gives the same sensitivity at a similar false-alarm rate, so the gain comes from whitening, not from calibration.

Over an eight-pulse dwell (SCR table in the Supplementary Material), PAMF-H is the most sensitive detector, and AR non-coherent integration is close to it. The ANMF (Tyler) needs at least 5.4 dB [1.9, 8.8] more SCR than PAMF-H; this is a lower bound, because PAMF-H reaches $P_d = 0.5$ at the lowest SCR tested. The P-ANMF pays for normalizing by a dwell that contains the target. The clipped integrator costs 2.0 dB [0.0, 2.9] against AR(1) integration at $\alpha = 0.01$, and the binary integrator costs more. Two classical alternatives are less sensitive. The power OS detector on the cell's own history (clutter-map OS-CFAR) needs 10.6 dB [7.0, 12.7] more SCR than IN-ARCP at the first look, so the sensitivity comes from whitening, not from the order statistic. A Hann-windowed eight-pulse Doppler filter bank with range CA normalization needs 5.0 dB more SCR than integrated IN-ARCP. All whitening detectors lose sensitivity for targets at the clutter Doppler (PAMF-H by 4.3 dB).

Real onsets come from the 77 GHz pedestrian run. Its dense interference is suppressed by fast-time zeroing. Then one chirp per frame (200 ms apart) is tested against the 16




Fig. 3. Detection of Swerling-1 targets injected into IPIX test episodes ($m = 16$, conformal thresholds, per-look $P_{fa} = 0.01$). (a) Target in the tested pulse (74,980 episodes in total). NA-AR(4): noise-aware score, 8.0 dB less SCR than CA16. Circles: law of Corollary 2, model fitted on each unit's training third. Curves average P_d over units; the text averages per-unit gains. (b) Persistent 10 dB target (56 units) without and with an eight-pulse guard, markers offset for visibility; means in the text are over looks 1–8. Lines: law of Proposition 4 at the fitted AR(1) coefficient, random Doppler (exploratory).

earlier frames of its bin, or its range neighbors for range CA-CFAR, at false-alarm rates of 0.96–1.10 times the design 10^{-2} . Of 74 onsets of the person's direct and bistatic returns in 15-cm range bins, marked by the radar's own range-Doppler map, IN-ARCP detects 0.30 within three frames (mean over 16 receiver-chirp tests), the power clutter map 0.27, its OS form 0.23 and range CA-CFAR 0.07. Frame-to-frame clutter is nearly uncorrelated here (median $|r| = 0.12$), so Proposition 2 predicts little whitening gain, and IN-ARCP only matches the clutter map. The onsets fall in 7 distinct five-frame blocks, fewer than the eight required in advance, so this result is descriptive (Supplementary Material).

C. Persistent and emerging targets

A target that stays in the beam after its onset enters the history. As Proposition 4 predicts, IN-ARCP then loses it (Fig. 3b): at 10 dB its per-look P_d falls from 0.885 at onset to 0.084 at look 4. A target at the clutter Doppler is already lost at look 1 (0.024), whereas one at the opposite Doppler is still detected there (0.940) and lasts several looks, in line with the visibility horizon of 3.7 looks (Corollary 4). The exact law reproduces these decays, with mean absolute errors of 0.010, 0.003 and 0.008 for random, clutter and opposite Doppler (exploratory).

A guard of eight pulses keeps the target out of the scale. With it, the per-look P_d over the eight guarded looks averages 0.82 for random and 0.95 for opposite Doppler, against 0.24 without the guard for random Doppler. The law gives 0.85 (exploratory), and the conformal false-alarm rate stays at 1.03 times design. The guard costs 0.9 dB at onset, plausibly because its scale is eight pulses older. It cannot help at the clutter Doppler, where whitening removes the target (0.07). The median OS score also keeps the target (P_d 0.650–0.801), but it costs 2.0 dB for abrupt targets, above the 0.69 dB predicted by Corollary 3; the guard is cheaper.

Targets that emerge gradually are harder for the per-look screen. When the target amplitude rises linearly over L pulses, reaching full amplitude at the first look, IN-ARCP's first-look gain over CAloc falls from 9.9 dB for an abrupt onset in this comparison to 4.7 dB for $L = 8$. At $L = 16$ the target is barely detected ($P_d = 0.014$ at 10 dB), whereas PAMF-H over the dwell loses only 1.9 dB. The exact law

TABLE II
REMEDIES FOR PULSED INTERFERENCE ON IPIX ($\alpha = 0.01$, 56 UNITS).
INTERFERENCE IS INJECTED 30 dB ABOVE THE LOCAL POWER; EACH PULSE IS HIT WITH PROBABILITY p_h , AND BURSTS ARE EIGHT PULSES LONG AT $p_h = 5\%$. ENTRIES ARE THE MEASURED P_{fa}/α AND THE P_d OF A 10 dB PERSISTENT TARGET. TRIMMED: 4 (2%) AND 5 (5%) OF EIGHT TERMS DISCARDED. BOLD: $P_{fa}/\alpha > 1.2$.

Detector	None	$p_h = 2\%$		$p_h = 5\%$		Bursts
	P_d	P_{fa}/α	P_d	P_{fa}/α	P_d	P_{fa}/α
Clipped, clean threshold	0.86	1.10	0.85	1.28	0.85	3.98
Certified clipped, $\kappa = 6$	–	0.55	0.81	0.12	0.64	1.25
Certified clipped, $\kappa = 9$	–	0.53	0.81	0.14	0.71	1.66
Certified trimmed	–	0.57	0.79	0.23	0.72	3.60
Blanking, mask-matched	0.79	1.08	0.79	1.21	0.78	4.14

reproduces the ramps (mean absolute error 0.016). The ramp fills the normalizing history, and with the eight-pulse guard the probability of detecting a 16-pulse ramp within eight looks at 10 dB rises from 0.03 to 0.87.

The IPIX reference target tests the opposite extreme. It is present throughout each recording, so it lies in every history and no guard helps. This blindness was seen during development and is explained by Proposition 3; a pre-specified test confirms it for IN-ARCP and the clipped integrator: both exceed their thresholds on the target cell at a rate of 0.0004 at $\alpha = 10^{-3}$, below their clutter false-alarm rates. Self-normalized whitened Doppler detects the target with $P_d = 0.22$ over 8 pulses and 0.40 over 256 pulses. Its measured clutter false-alarm rates there are 1.1 and 1.5×10^{-3} , above the design 10^{-3} (Supplementary Material). Range CA integration barely detects the target ($P_d \leq 0.007$), and the ANMF reaches 0.13–0.18 over 8–16 pulses. Learned detectors report $P_d = 0.91$ at $P_{fa} = 10^{-3}$ on IPIX with 1.024-s observations and a different evaluation [16].

D. Pulsed interference

With injected pulsed interference on IPIX at $\alpha = 10^{-2}$, non-coherent integration (AR(1) or AR(4)) and PAMF-H exceed their design false-alarm probability 10- to 23-fold over two hit rates and two powers, because a dwell pulse corrupted by interference is an outlier the history does not predict. Self-normalized detectors (P-ANMF, ANMF) stay within 1.21 times design but lose sensitivity.

The clip level matters at low false-alarm rates. Without interference at 10^{-3} , $\kappa = 6$ makes the clipped integrator conservative through saturation ($P_{fa}/\alpha = 0.67$). A pre-specified saturation rule on calibration data most often picks $\kappa = 9$ (17 of 56 units) and is not met by any tested clip level ($\kappa \leq 18$) in 16 units. Applied to all units (post hoc), $\kappa = 9$ gives $P_{fa}/\alpha = 0.90$ and raises P_d at 10 dB from 0.46 to 0.68, with negligible loss at 10^{-2} .

Table II compares the remedies with dense calibration (about 14,600 dwell segments per unit). The certified clipped integrator stays below design, as Theorem 1 guarantees. The injected hits follow the calibration hit model, so this checks the implementation and the cost, not robustness. Its P_d at 10 dB falls from 0.86 with a clean threshold and no interference to 0.81 at a 2% hit rate and 0.64 at 5% (0.71 with $\kappa = 9$). The certified trimmed sum performs similarly, and so does certified binary integration (P_d within 0.03; exploratory, Supplementary Material). At 10^{-3} with 5% hits, the certificate is vacuous for every clip level: the certified threshold detects no target.

Pulse blanking, calibrated on clean episodes blanked at hit positions from the hit model, is cheaper but not certified. It keeps $P_d = 0.78$ at 5% hits and stays within 1.21 times design at 10^{-2} . At 10^{-3} with 5% hits it exceeds design 1.7- to 2.2-fold. Bursts of eight pulses at the same hit rate defeat every remedy (Table II): a certificate calibrated on independent hits exceeds design, and one calibrated on bursts never detects.

On the 77 GHz data the interference is real and occupies whole chirps. The hit model draws random 24-chirp blocks from the hit masks of frames 0–49; the test uses frames 50–99. In the sparse run the certified clipped integrator holds the false-alarm rate at 0.54 times design without fast-time mitigation, whereas non-coherent integration runs at 1.42 times design overall and 3.56 times in segments with one or two hit chirps. Against clipping with a clean threshold, the certificate costs 3.1 dB on hit-free segments and 6.4 dB overall. In the dense run, real interference mainly masks targets rather than raising false alarms (non-coherent integration 0.81 times design), and certification is nearly vacuous. Fast-time zeroing returns the uncertified detectors to their clean-reference false-alarm rates.

E. A second sea-clutter radar

Table III repeats the key tests on NetRAD, unused in development. At 1 kHz, NetRAD clutter is more correlated than IPIX clutter ($|r| = 0.97$ –0.98, at or near the 0.98 clamp, against an IPIX median of 0.96), so whitening gains more: IN-ARCP's first-look gain over the raw-power OS detector is ≥ 12 dB (Supplementary Material). The model-based laws again fail in the tail, and the guard reproduces the IPIX result: the per-look P_d is 0.92 with the guard against 0.30 without it (law 0.92, exploratory), and the probability of detecting a 16-pulse ramp within eight looks rises from 0.02 to 0.94.

Calibration transfers at 10^{-3} but not fully at 10^{-4} . Conformal thresholds hold all thirteen statistics, the clipped integrator included, within 1.01–1.36 times design at 10^{-3} , and the IN-ARCP, AR(4), OS and power scores within 1.38 at 10^{-4} . At 10^{-4} , however, five statistics exceed twice design: the noise-aware score (2.08), integration, PAMF-H and the two P-ANMF statistics (up to 4.26). An exploratory analysis traces

TABLE III
IPIX AND NetRAD S-BAND SEA CLUTTER, SAME CODE AND SETTINGS EXCEPT THE ANMF SECONDARY CELLS (SECTION V). ENTRIES: P_{fa}/α AT $\alpha = 0.01$ UNLESS STATED; GUARD ROW: PER-LOOK P_d AT 10 dB, MEAN OVER LOOKS 1–8; P_d IN PARENTHESES AT 10 dB. CONFORMAL ROWS: THE 12 STATISTICS OTHER THAN THE CLIPPED INTEGRATOR (CONSERVATIVE ON IPIX). INJECTED INTERFERENCE 30 dB ABOVE THE LOCAL POWER. FLAGGED PULSES: NetRAD HH, CERTIFICATE CALIBRATED ON THEIR MASKS.

	IPIX	NetRAD
Conformal, 10^{-3} (12 statistics)	0.98–1.14	1.01–1.36
Conformal, 10^{-4} (12 statistics)	0.73–1.47	1.02–4.26
Model-based laws $> 2\times$ at 10^{-4}	10 of 11	9 of 11
IN-ARCP CA law at 10^{-4}	1.92	1.70
Guard: P_d per look, $\Delta = 0 / 8$	0.24 / 0.82	0.30 / 0.92
5% hits: AR(1) integration	21	22
5% hits: certified clip (P_d)	0.12 (0.64)	0.08 (0.92)
5% bursts: blanking	4.14	4.11
Flagged pulses: integration / certified (P_d)	–	1.05 / 0.04 (0.04)

most of the integration and PAMF-H excess to two units whose exceedances coincide with flagged pulses, while their calibration thirds contain none (without them, 1.56 and 1.33). The P-ANMF excess comes from one unit with flagged pulses in both thirds, and fewer than half of its exceedances coincide with them (exploratory). IN-ARCP's texture spread is 2.40, larger than CA16's (1.51), whereas the AR(4) score's is 1.11, so at the clamp AR(1) whitening is plausibly incomplete.

The rule flags 3–23% of the HH pulses, in bursts of several pulses, yet uncertified non-coherent integration stays at 1.05 times design at 10^{-2} (1.20 at 10^{-3}). A certificate calibrated on these masks is nearly vacuous ($P_d = 0.04$), plausibly because the flagged pulses are frequent and come in bursts.

VII. DISCUSSION

A. Where the detector works and where it fails

In Table IV, the failures have four sources: a target that enters the normalizing scale (persistent without a guard, a ramp, the IPIX target); interference that enters the dwell of an uncertified detector; a certificate whose hit model must cover bursts or frequent hits, at a cost that grows to vacuity; and a threshold taken from a model rather than from data. A persistent target at the clutter Doppler is lost after onset even with a guard, and data-calibrated thresholds still depart from design on NetRAD at 10^{-4} .

The results also rest on conditions that were assumed or only partly tested. The laws assume the true AR model and pure compound-Gaussian clutter; on real clutter they predict detection closely, but not the far tail. Conformal validity is marginal and needs exchangeable calibration and test episodes, which holds only approximately within a session; online recalibration is vulnerable to target contamination (Supplement). The certificate needs hit positions independent of the clutter and dominated by the hit model, and interference too weak to be flagged counts as clutter. Apart from the IPIX reference target and the walking person, whose result is descriptive, the real-data detection results rest on injected targets with known timing, to which the dwell is aligned. A guard extends the horizon by only Δ looks. The sea-clutter evidence comes from

TABLE IV

WHERE THE DETECTOR WORKS AND WHERE IT FAILS, OVER THE SCENARIOS TESTED IN SECTION VI. W: MEETS DESIGN OR THE LAW; P: MEETS DESIGN AT A COST, OR WITH A CAVEAT; F: ABOVE TWICE DESIGN, OR NO DETECTION. A SLASH SEPARATES THE TWO CONDITIONS OF A ROW; GRADES FOLLOW THE WORST CASE. P_{fa}/α AT $\alpha = 10^{-2}$ AND P_d AT 10 dB UNLESS STATED.

Scenario	Result	Scenario	Result
<i>False-alarm control</i>		<i>Targets (continued)</i>	
Conformal, IPIX, to 10^{-4}	W 12 of 13 in 0.73–1.47	At the clutter Doppler	F 0.024 at look 1; 0.07 with guard
Conformal, NetRAD, $10^{-3} / 10^{-4}$	W / F $\leq 1.36 / 5$ of 13 above $2\times$	In every history (IPIX target)	F / P IN-ARCP blind; P-ANMF 0.22–0.40 at 10^{-3}
Model laws, 10^{-4} , IPIX / NetRAD	F 10 / 9 of 11 above $2\times$	<i>Pulsed interference</i>	
Across clutter power, IPIX	W spread 1.7 (raw RMS 9.8)	Uncertified integration, 2–5%	F 10- to 23-fold
Across CNR, 77 GHz	P 1.07 (CNR < -5 dB) to 0.08 (10–20 dB)	Certified, 2–5% / 5% at 10^{-3}	P / F $\leq \alpha$, P_d 0.81 / 0.64 ($\kappa = 6$) / no detection
<i>Targets</i>		Blanking, isolated, $10^{-2} / 10^{-3}$	P / F $\leq 1.21 / 2.2$ -fold at 5%
Abrupt onset	W 7.0 dB over CA16; law error 0.021	Bursts, every remedy	F cert. 1.66-fold or vacuous, blanking 4.1-fold
Persistent, guard $\Delta = 0 / 8$	F / W 0.24 (0.084 at look 4) / 0.82	Real 77 GHz, sparse / dense	P / F 0.54 at a cost / nearly vacuous
16-pulse ramp, within 8 looks	F / W 0.03 / 0.87 without / with guard	Fast-time zeroing, 77 GHz	W clean-reference rates
Real onsets, 3 frames, 77 GHz	P 0.30 vs 0.27 clutter map, descriptive	NetRAD flagged pulses, certified	F P_d 0.04

two radars with one campaign each, and the real onsets from one walk, tracked for about six seconds, in 77 GHz ground clutter, tested at frame rate and marked by the radar's own range-Doppler processing, without an external sensor.

B. Design rules for a cell-wise detector

- 1) Whiten with an AR model fitted on clutter training data, and normalize by the RMS innovation scale over a window that ends Δ pulses before the tested pulse; the guard keeps a target out of the scale for Δ looks ($\Delta = 8$ on both radars). Targets present throughout the history need self-normalized statistics (modest P_d on the real IPIX target).
- 2) Calibrate thresholds conformally on clean clutter: beyond about $3.15/\alpha$ episodes, $\mathbb{P}\{P_{fa} > 2\alpha\} \leq 0.05$, and more are needed under dependence.
- 3) Excise pulsed interference in fast time where possible. Otherwise, at low hit rates, integrate clipped innovations ($\kappa \approx 9$) with a certified threshold and a hit model that dominates the hit pattern, not only its rate. Neither certification nor blanking handles bursts.

VIII. CONCLUSION

At the true coefficient, whitening a range cell with an AR model and normalizing by its own innovation power turns correlated compound-Gaussian clutter with AR speckle and no thermal noise into the independent exponential samples of CFAR theory, so the classical laws hold with the whitened SCR. Further laws follow: a post-onset law whose visibility horizon tells when a persistent target masks itself, a guard that extends this horizon by its length, the OS counterpart, and a certificate for clipped integration under pulsed interference of any power. On sea clutter of two radars, model-based thresholds fail in the tail, whereas conformal calibration keeps every statistic at or below 1.36 times design at 10^{-3} , and 12 of 13 within 1.47 at 10^{-4} on IPIX, though not on NetRAD. An eight-pulse guard keeps an injected persistent 10 dB target visible over the eight guarded looks (per-look P_d 0.82 on IPIX, 0.92 on NetRAD), except at the clutter Doppler, and on IPIX the law at onset predicts detection probability. The certificate holds on sea clutter and on real FMCW interference at a cost that grows to vacuity at high hit rates; fast-time excision avoids it where it applies. Bursty interference remains

open; hit models that count only harmful hits, and real-target campaigns at sea, are the next steps.

APPENDIX A PROOF OF PROPOSITION 1

Write $X = \Sigma^{1/2}U$ with $U \sim \mathcal{CN}(0, I)$. The event is $\{X^H M_q X \leq 0\}$, and $X^H M_q X = \sum_k \lambda_k E_k$ with independent standard exponentials E_k . M_q is positive in the direction e_{m+1} , and $M_q + \frac{q^2}{m} \text{diag}(Q, 0)$ has rank one, so by Weyl's inequality M_q has exactly one positive eigenvalue; by Sylvester's law of inertia so does $\Sigma^{1/2} M_q \Sigma^{1/2}$. All other nonzero λ_k are negative. Conditioning on $Z = \sum_{k \neq +} |\lambda_k| E_k$ and using $\mathbb{P}\{E_+ > z\} = e^{-z}$,

$$\mathbb{P}\{\lambda_+ E_+ > Z\} = \mathbb{E} e^{-Z/\lambda_+} = \prod_{k \neq +} \frac{\lambda_+}{\lambda_+ - \lambda_k},$$

whatever the multiplicities, and (2) is its complement.

APPENDIX B PROOFS OF COROLLARY 2, PROPOSITIONS 3 AND 4, AND THEOREM 1

Corollary 2. At $r = \rho$ in pure AR(1) clutter of power σ^2 , $Y - \rho H_m = \varepsilon + \sqrt{S} \sigma g$ with independent $\varepsilon \sim \mathcal{CN}(0, \sigma^2(1 - |\rho|^2))$, $g \sim \mathcal{CN}(0, 1)$, and $ms_\rho^2/(\sigma^2(1 - |\rho|^2))$ is a Gamma(m) variable; (3) is its moment-generating function.

Proposition 3. The score is homogeneous of degree zero, so its value at $C + Au$, with clutter C and $u = (e^{j\omega t})_t$, equals that at $C/A + u$ and tends to its value at u ; at u the numerator is $|1 - re^{-j\omega}|$ and $ms_r(u)^2 = (1 - |r|^2) + (m-1)|1 - re^{-j\omega}|^2$. Dominated convergence gives $P_d \rightarrow 0$.

Proposition 4. In whitened units the episode is $\mathcal{CN}(0, I + vv^H)$, with $|v_{m+1}|^2 = \mathcal{E}_0$ and $\sum_{t \leq m} |v_t|^2 = \mathcal{E}_H$, and the event is $x^H M x > 0$, $M = \text{diag}(-\beta I_m, 1)$. The eigenvalues of $M(I + vv^H)$ are $-\beta$ ($m-1$ times) and the roots of $(\lambda-1)(\lambda+\beta) - \mathcal{E}_0(\lambda+\beta) + \beta \mathcal{E}_H(\lambda-1) = 0$ (rank-one update of a diagonal matrix). With $\lambda_+ > 0 > \lambda_-$, $\mathbb{P}\{\lambda_+ E_+ > |\lambda_-| E_- + \beta \Gamma_{m-1}\}$, with Γ_{m-1} a sum of $m-1$ unit exponentials, is (5). For the horizon, $\lambda_+ \rightarrow \infty$ or stays bounded as $S \rightarrow \infty$ according to the sign of $|d|^2 - \beta\{1 + (\ell-1)|d|^2\}$.

Theorem 1. The k_H -th smallest history power is nondecreasing in each power, so for any values of the corrupted entries it is at least its value with those entries set to zero. Uncorrupted

dwell terms are then clean powers over a scale at least that large, and corrupted clipped terms are at most κ . Hence $T \leq W(x, \mathcal{J})$, and W is nondecreasing in $\mathcal{J}$. If $\mathcal{J} \subseteq \tilde{\mathcal{J}} \sim \mathcal{D}$ with $\tilde{\mathcal{J}}$ independent of the clutter, $T \leq W(X, \tilde{\mathcal{J}})$, and $(X_i, \mathcal{J}_i)$ and $(X, \tilde{\mathcal{J}})$ are exchangeable, so the split-conformal rank argument bounds the false-alarm probability by α . Details and the remaining laws are in the Supplementary Material.

CODE AND DATA AVAILABILITY

Code, hashed protocols, outcomes, figure scripts and the package `inarc` 0.3.0: <https://github.com/razaumair2203-ux/inarc>, release tag `r17-trs-submission`. Data: IPIX from McMaster University [37], NetRAD from University College London [39], 77 GHz from Zenodo [38].

ACKNOWLEDGMENT

The authors thank the McMaster University IPIX group, M. Ritchie and the UCL and University of Cape Town NetRAD team, and L. Rienessl and R. Feger (JKU Linz) for their open radar data. Anthropic Claude Code assisted with the derivations, simulations, analysis code and the preparation of the text in all sections; OpenAI Codex assisted earlier versions with code and editing. The authors designed and validated the technique, the experiments and the test plan, and take full responsibility for the content.

REFERENCES

- [1] H. M. Finn and R. S. Johnson, "Adaptive detection mode with threshold control as a function of spatially sampled clutter-level estimates," *RCA Review*, vol. 29, no. 3, pp. 414–464, 1968.
- [2] H. Rohling, "Radar CFAR thresholding in clutter and multiple target situations," *IEEE Trans. Aerosp. Electron. Syst.*, vol. AES-19, no. 4, pp. 608–621, 1983.
- [3] P. Gandhi and S. Kassam, "Analysis of CFAR processors in nonhomogeneous background," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 24, no. 4, pp. 427–445, 1988.
- [4] M. A. Richards, *Fundamentals of Radar Signal Processing*, 2nd ed. New York: McGraw-Hill Education, 2014.
- [5] F. S. Abbasi, U. I. Ahmed, and S. Ahmed, "Combating radar pulse jamming using clipping-based noncoherent pulse integration," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 54, no. 4, pp. 1783–1789, Aug. 2018.
- [6] K. Ward, R. Tough, and S. Watts, *Sea Clutter: Scattering, the K Distribution and Radar Performance*. Institution of Engineering and Technology, 2013.
- [7] E. Conte, M. Lops, and G. Ricci, "Asymptotically optimum radar detection in compound-Gaussian clutter," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 31, no. 2, pp. 617–625, 1995.
- [8] E. Ollila, D. E. Tyler, V. Koivunen, and H. V. Poor, "Complex elliptically symmetric distributions: Survey, new results and applications," *IEEE Trans. Signal Process.*, vol. 60, no. 11, pp. 5597–5625, 2012.
- [9] E. J. Kelly, "An adaptive detection algorithm," *IEEE Trans. Aerosp. Electron. Syst.*, vol. AES-22, no. 2, pp. 115–127, Mar. 1986.
- [10] S. Kraut and L. Scharf, "The CFAR adaptive subspace detector is a scale-invariant GLRT," *IEEE Trans. Signal Process.*, vol. 47, no. 9, pp. 2538–2541, 1999.
- [11] F. Pascal, Y. Chitour, J.-P. Ovarlez, P. Forster, and P. Larzabal, "Covariance structure maximum-likelihood estimates in compound Gaussian noise: Existence and algorithm analysis," *IEEE Trans. Signal Process.*, vol. 56, no. 1, pp. 34–48, Jan. 2008.
- [12] J. Roman, M. Rangaswamy, D. Davis, Q. Zhang, B. Himed, and J. Michels, "Parametric adaptive matched filter for airborne radar applications," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 36, no. 2, pp. 677–692, 2000.
- [13] J. Michels, B. Himed, and M. Rangaswamy, "Evaluation of the normalized parametric adaptive matched filter STAP test in airborne radar clutter," in *Record of the IEEE 2000 International Radar Conference*. IEEE, 2000, pp. 769–774.
- [14] T. Bai, J. Tang, and Y. Xu, "Complex VAE with heavy-tailed likelihood for radar target detection in sea clutter," arXiv:2606.10540, 2026.
- [15] Y. A. Rouzoumka, J. Pinsolle, E. Terreaux, C. Morisseau, J.-P. Ovarlez, and C. Ren, "Out-of-distribution radar detection with complex VAEs: Theory, whitening, and ANMF fusion," arXiv:2601.18677, 2026.
- [16] Q. Qu, W. Liu, J. Wang, B. Li, N. Liu, and Y.-L. Wang, "Enhanced CNN-based small target detection in sea clutter with controllable false alarm," *IEEE Sensors J.*, vol. 23, no. 9, pp. 10 193–10 205, May 2023.
- [17] C. Qu, X. Wang, and J. Yin, "Sea clutter radar target detector based on autoregressive sieve bootstrap," *Digit. Signal Process.*, vol. 145, p. 104298, 2024.
- [18] C. Qu, L. Yang, L. Zhen, X. Wang, and J. Yin, "Robust adaptive CFAR detection using AR-sieve bootstrap and sample autocovariance," *Digit. Signal Process.*, vol. 168, p. 105574, Jan. 2026, art. no. 105574.
- [19] M. Lops and M. Orsini, "Scan-by-scan averaging CFAR," *IEE Proc. F*, vol. 136, no. 6, pp. 249–254, 1989.
- [20] M. Naldi, "False alarm control and self-masking avoidance by a biparametric clutter map in a mixed interference environment," *IEE Proc.-Radar Sonar Navig.*, vol. 146, no. 4, pp. 195–200, 1999.
- [21] J. Lei, M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, "Distribution-free predictive inference for regression," *J. Amer. Statist. Assoc.*, vol. 113, no. 523, pp. 1094–1111, 2018.
- [22] V. Vovk, "Conditional validity of inductive conformal predictors," in *Proc. ACML*, vol. 25, 2012, pp. 475–490.
- [23] S. Haykin and X. B. Li, "Detection of signals in chaos," *Proc. IEEE*, vol. 83, no. 1, pp. 95–122, 1995.
- [24] S. Bates, E. Candès, L. Lei, Y. Romano, and M. Sesia, "Testing for outliers with conformal p-values," *Ann. Statist.*, vol. 51, no. 1, 2023.
- [25] S. D. Himonas, "CFAR integration processors in randomly arriving impulse interference," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 30, no. 3, pp. 809–817, Jul. 1994.
- [26] M. Toth, P. Meissner, A. Melzer, and K. Witrals, "Performance comparison of mutual automotive radar interference mitigation algorithms," in *Proc. 2019 IEEE Radar Conference (RadarConf)*. IEEE, Apr. 2019, pp. 1–6.
- [27] M. Rameez, M. Dahl, and M. I. Pettersson, "Autoregressive model-based signal reconstruction for automotive radar interference mitigation," *IEEE Sensors J.*, vol. 21, no. 5, pp. 6575–6586, Mar. 2021.
- [28] A. Gendler, T.-W. Weng, L. Daniel, and Y. Romano, "Adversarially robust conformal prediction," in *International Conference on Learning Representations (ICLR)*, 2022.
- [29] S. H. Zargarbashi, M. S. Akhondzadeh, and A. Bojchevski, "Robust yet efficient conformal prediction sets," in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, vol. 235. PMLR, 2024, pp. 17 123–17 147.
- [30] A. Levine and S. Feizi, "Robustness certificates for sparse adversarial attacks by randomized ablation," *Proc. AAAI Conf. Artif. Intell.*, vol. 34, no. 4, pp. 4585–4593, 2020.
- [31] M. Zaffran, A. Dieuleveut, J. Josse, and Y. Romano, "Conformal prediction with missing values," in *Proc. 40th International Conference on Machine Learning (ICML)*, vol. 202, 2023, pp. 40 578–40 604, arXiv:2306.02732.
- [32] S. Watts, "Radar detection prediction in K-distributed sea clutter and thermal noise," *IEEE Trans. Aerosp. Electron. Syst.*, vol. AES-23, no. 1, pp. 40–45, 1987.
- [33] N. Ignatiadis and B. Sen, "Empirical partially Bayes multiple testing and compound χ^2 decisions," *Ann. Statist.*, vol. 53, no. 1, 2025.
- [34] J.-P. Imhof, "Computing the distribution of quadratic forms in normal variables," *Biometrika*, vol. 48, no. 3–4, pp. 419–426, 1961.
- [35] F. Gini, "Sub-optimum coherent radar detection in a mixture of K-distributed and Gaussian clutter," *IEE Proc.-Radar Sonar Navig.*, vol. 144, no. 1, pp. 39–48, 1997.
- [36] R. A. Fisher, "Tests of significance in harmonic analysis," *Proc. R. Soc. Lond. A*, vol. 125, no. 796, pp. 54–59, Aug. 1929.
- [37] McMaster University IPIX Radar Group, "The IPIX radar database: Dartmouth 1993," <https://soma.ece.mcmaster.ca/ipix/dartmouth/data-sets.html>, accessed 29 Sep. 2026.
- [38] L. Rienessl and R. Feger, "77-GHz FMCW radar network ADC data under various interference scenarios at parking lot," Zenodo dataset, doi:10.5281/zenodo.21933926, 2026.
- [39] M. Ritchie, "NetRAD – monostatic & bistatic sea clutter dataset," University College London, dataset, doi:10.5522/04/32676582, 2026.
- [40] F. Fioranelli, M. Ritchie, H. Griffiths, S. Sandenbergh, and M. Inggs, "Analysis of polarimetric bistatic sea clutter using the NetRAD radar system," *IET Radar Sonar Navig.*, vol. 10, pp. 1356–1366, 2016.