

Innovation-Normalized Detection in Compound-Gaussian Clutter: Exact Laws, Conformal Thresholds and Certified Integration under Pulsed Interference

Muhammad Umair Raza, Sohail Ahmed, Ammad Ahmed, M. Atif Shahzad, and Syed M. Kazam Abbas Kazmi

Abstract—Classical cell-averaging and order-statistic constant-false-alarm-rate (CFAR) laws assume independent exponential samples. Correlated, textured clutter can violate this premise. Autoregressive (AR) whitening and innovation-power normalization recover these laws under the true AR compound-Gaussian model with common episode texture and no thermal noise. We derive a post-onset law quantifying persistent-target contamination of a cell's own history; a slow-time guard delays self-masking by its length. Under exchangeability and a dominating, clutter-independent hit model, a conformal threshold on a worst-case bound certifies clipped and binary integration against pulsed interference of any power. Simulations and measurements of sea clutter and 77 GHz ground scenes test calibration, target visibility, noise and interference. On McMaster IPIX sea clutter, conformal thresholds keep 12 of 13 statistics within 0.73–1.47 times design down to 10^{-4} ; 10 of 11 model-based laws exceed twice design, up to 60-fold. At 10^{-2} , an eight-pulse guard raises mean per-look detection of an injected persistent 10 dB target over looks 1–8 from 0.24 to 0.82, except at clutter Doppler; 48-look trajectories show later collapse. The guard benefit recurs on NetRAD S-band clutter, unused in initial development; five statistics exceed twice design at 10^{-4} after threshold refitting. With sparse real 77 GHz interference at 10^{-2} , certified clipping gives 0.54 times design and costs 6.4 dB versus clean-threshold clipping; fast-time zeroing restores clean-reference rates. Radar- and camera-labelled pedestrians show little whitening advantage at weak inter-frame correlation. Guarded acquisition remains poor under idealized continuous migration; burst-matched certification controls false alarms but detects no target.

Index Terms—Compound-Gaussian clutter, conformal prediction, constant false alarm rate, pulse integration, pulsed interference, sea clutter.

I. INTRODUCTION

CLASSICAL CFAR design uses exact false-alarm and detection laws for cell-averaging (CA), order-statistic (OS), non-coherent, binary and clipped integration, under independent exponential powers and specified target models [1]–[5]. These assumptions can fail in compound-Gaussian clutter: texture scales correlated Gaussian speckle, and thermal noise adds a floor [6]–[8]. Exponential-law thresholds can then lose both power invariance and tail accuracy.

M. U. Raza, M. A. Shahzad, and S. M. K. A. Kazmi are with the Department of Avionics Engineering, College of Aeronautical Engineering, National University of Sciences and Technology (NUST), Risalpur, Pakistan (corresponding author: M. U. Raza; e-mail: uraza@cae.nust.edu.pk).

S. Ahmed is with MathWorks, Natick, MA, USA.

A. Ahmed is with the National Disaster Management Authority (NDMA), Islamabad, Pakistan.

A detector that tests each range cell along slow time faces three further problems. First, its threshold must hold far in the tail. Adaptive covariance and AR-whitening detectors restore scale invariance under their model assumptions (Section II). This does not establish the accuracy of model-based thresholds in the far tail of measured clutter. Learned and resampling sea-clutter detectors set thresholds empirically [9]–[13], whose guarantees and performance models differ from the rank and detection laws developed here. Second, a detector that normalizes a cell by its own past can mask a target that persists into that past, as clutter-map detectors do [14], [15]. Third, pulsed interference from other radars, jammers or the chirp crossings of frequency-modulated continuous-wave (FMCW) radars corrupts a few pulses with unknown power, and with them both the test statistic and its reference.

We address these problems by AR-whitening each cell's slow-time history with a clutter-trained model and normalizing the tested innovation by the history innovation power (Fig. 1). At the true AR coefficient and without thermal noise, the innovation powers are conditionally independent exponentials; their common texture scale cancels in the normalization, so the classical laws return with the signal-to-clutter ratio (SCR) multiplied by the whitening factor. The per-look score is a one-pulse parametric adaptive matched filter (PAMF) test [16]. We call it IN-ARCP (innovation-normalized autoregressive conformal prediction), because its exceedance region is the complement of a conformal prediction disc for the next sample [17]. The contributions are as follows.

- 1) **Detection after onset** (Section IV). Building on classical CA, OS and integration laws for innovations, we derive a post-onset rank-one detection law and its OS counterpart. It gives a *visibility horizon* for strong persistent-target self-masking. A slow-time guard, analogous to range-CFAR guards [4], delays this horizon by its length. Under the model, all 65 Monte Carlo checks lie within 3 standard errors (Supplementary Material).
- 2) **Calibration** (Section IV-F). The split-conformal finite-sample law [18] sets the calibration size. On IPIX and NetRAD, these thresholds improve measured tail calibration over the tested model-based laws, including those of the adaptive normalized matched filter (ANMF); NetRAD retains failures at 10^{-4} .
- 3) **Certified integration under pulsed interference** (Theo-

rem 1). For any exchangeable clean-episode law, a conformal threshold on a worst-case bound for clipped [5] and binary integration bounds the false-alarm probability for interference of any power. The pattern of interfered pulses must be independent of the clutter and dominated, in the inclusion order, by the hit model used in calibration.

We evaluate tail calibration on X-band IPIX and S-band NetRAD sea clutter; NetRAD was unused in initial development and its thresholds are refitted. Injected onsets, longer trajectories and continuous migration test sensitivity to reference contamination. The 77 GHz Johannes Kepler University (JKU) recordings test real interference and pedestrian returns; University of Washington (UW) recordings add detector-external camera labels. Comparators include PAMF, ANMF, power detectors, certified trimming and pulse blanking.

II. RELATED WORK

Normalized detection in compound-Gaussian clutter. Scale invariance of normalized statistics underlies texture-CFAR detection [7], [8], [19]. The ANMF, also called the adaptive coherence estimator (ACE), uses covariance estimated from secondary data: the sample covariance matrix (SCM) [19] or the fixed-point (Tyler) estimator [20]; the PAMF and normalized PAMF whiten with an AR filter [16], [21]. Detection through the prediction error of a clutter model is classical, including on IPIX [22]. Targets in the reference window mask CA detectors, and they mask OS detectors once they occupy more than $m - k$ reference cells [2], [3]. Clutter-map detectors filter a cell's own past power [14], and their self-masking by persistent targets has been analyzed and countered by censoring [15]. We found no such analysis for the AR-innovation score of a cell's own history. Motion-aware coherent integration searches jointly over range and velocity [23]. Our migration study tests whether the cell-wise history-normalized score remains useful when crossing time is unknown.

Data-driven thresholds. Split conformal prediction gives finite-sample marginal coverage for exchangeable data [17], [18], and a conformal p -value is a calibrated test [24]. Conformal calibration has also controlled false alarms in SAR change detection [25]. Recent learned radar detectors use empirical null distributions with concentration bounds [10]. Rank-based nonparametric CFAR also obtains distribution-free null laws from i.i.d. continuous range-reference samples [26]. Here split conformal ranks complete episode scores under exchangeability; innovation normalization makes the score nearly pivotal where clutter dominates thermal noise.

Integration under pulsed interference. Non-coherent, binary and censored integration in random impulse interference was analyzed in [27], and clipping the per-pulse statistics is a robust remedy [5]; for automotive FMCW interference, fast-time zeroing and reconstruction are standard [28], [29]. Robust conformal prediction certifies coverage under bounded or sparse perturbations [30], [31], randomized ablation certifies against sparse perturbations of unbounded amplitude [32], and calibrating over masks that contain the test mask handles missing values [33]. Theorem 1 applies the mask device to

pulsed interference, where order statistics give a closed-form bound over all amplitudes.

III. MODEL AND DETECTORS

A. Compound-Gaussian autoregressive episodes with thermal noise

An *episode* $X = (H_1, \dots, H_m, Y) \in \mathbb{C}^{m+1}$ contains m consecutive samples of one range cell followed by the tested sample Y . Conditional on a latent texture,

$$X | c \sim \mathcal{CN}(0, \nu(cR + I)), \quad (1)$$

where R is the unit-diagonal speckle correlation, $\nu \geq 0$ the thermal-noise power and $c > 0$ the clutter-to-noise ratio (CNR), which carries the texture. Letting $\nu \rightarrow 0$ with νc fixed gives the compound-Gaussian model [7], [8]; the additive term is the thermal noise of K+noise clutter [6], [34]. A positive texture of arbitrary distribution scales stationary complex AR(p) speckle throughout each episode. The physical clutter source does not enter this model. For guarded or dwell tests, the shared texture covers the full scale and test span. Calibration validity assumes exchangeable episodes: their joint law is unchanged by reordering; disjoint times or cells do not ensure this. A *look* is one pulse test; $\ell = 0$ denotes onset and $\ell \geq 1$ the pulses elapsed since onset. A *dwell* is K consecutive tested pulses after the history.

B. Whitened innovations and detectors

Every whitened detector uses the innovations of an AR model fitted by least squares on training episodes. For AR(1) with coefficient r , radially clipped to $|r| \leq 0.98$ (the clip binds in 5 of 56 IPIX units), $\varepsilon_1 = \sqrt{1 - |r|^2}h_1$ and $\varepsilon_j = h_j - rh_{j-1}$; for AR(p), the innovations are the $m - p$ conditional ones.

- *Per-look score (IN-ARCP).* $|Y - rh_m|/s_r(h)$ with $s_r(h)^2 = m^{-1} \sum_j |\varepsilon_j|^2$. Its exceedance of the threshold $\hat{q}$ (Section III-C) is the complement of the prediction disc $\{y : |y - rh_m| \leq \hat{q}s_r(h)\}$. With a one-pulse steering vector, the PAMF statistic normalized by the history innovations equals, by construction, the squared AR(p) score (IN-AR(p)). With a *guard* of Δ pulses, the scale is computed from the m samples that end Δ pulses before h_m , while the center still uses h_m (Fig. 1); $\Delta = 0$ is IN-ARCP.
- *OS score.* The scale is the k -th smallest of the m history innovation powers.
- *Dwell detectors.* Over K dwell pulses after an m -pulse history, non-coherent integration sums the dwell innovation powers against their history mean. Clipped integration computes $\sum_{i=1}^K \min\{|\varepsilon_{m+i}|^2/s_{k_H}^2, \kappa\}$ with clip level κ , and binary integration counts the dwell powers $|\varepsilon_{m+i}|^2/s_{k_H}^2$ above η and thresholds the count; $s_{k_H}^2$ is the k_H -th smallest history innovation power. PAMF-H is the coherent PAMF over the dwell, normalized by the history innovations. Pulse blanking, a non-certified alternative, discards an innovation power and the one after it when the power exceeds both ten times the history median and four times a robust level of its block, and integrates the rest (Supplementary Material).




Fig. 1. Cell-wise detection after range processing. The scale uses $x_{t-\Delta-m}, \dots, x_{t-\Delta-1}$; the predictor retains the latest sample $h_m = x_{t-1}$, including when it lies in the guard. The AR coefficient is fitted on clutter training data and $\hat{q}$ on clean calibration episodes. The diagram shows the squared score and threshold.

- *Self-normalized whitened Doppler (parametric ANMF, P-ANMF).* $\max_j |\text{DFT}_j(\varepsilon)|^2 / (N \|\varepsilon\|^2)$ over N dwell innovations: the ANMF with an AR covariance, normalized by the dwell itself rather than by the history.
- *Noise-aware (NA) score.* The Gaussian predictive mean and variance of (1) given h , with (R, ν) and a nonparametric texture law fitted by marginal likelihood and a per-episode CNR [35]. Its ideal AR(1) zero-noise limit is IN-ARCP (Supplement).

C. Conformal thresholds

For a score T and a design false-alarm rate α , the threshold $\hat{q}$ is the k_α -th smallest value of T over n clean calibration episodes, with $k_\alpha = \lceil (n+1)(1-\alpha) \rceil$ and $\hat{q} = +\infty$ if $k_\alpha = n+1$. For exchangeable calibration and test episodes, the false-alarm probability is then at most α [17], [24]. Section IV-F gives its distribution given the calibration set.

IV. LAWS OF WHITENED-INNOVATION DETECTION

We first establish the onset law, then quantify history contamination and dwell integration before calibrating and certifying thresholds.

A. The exact law at onset

Proposition 1. For $q > 0$, let $X \sim \mathcal{CN}(0, \Sigma)$, $\Sigma \succ 0$, fix a linear center $a^\top H$ and a scale $s(H)^2 = H^\top Q H / m$ with $Q \succeq 0$, $Q \neq 0$, and let λ_k be the nonzero eigenvalues, with multiplicity, of $\Sigma^{1/2} M_q \Sigma^{1/2}$, $M_q = \bar{b} \bar{b}^\top - \frac{q^2}{m} \text{diag}(Q, 0)$, $b = e_{m+1} - (a^\top, 0)^\top$. Exactly one eigenvalue λ_+ is positive, and

$$G_\Sigma(q) := \mathbb{P}\{|Y - a^\top H| \leq q s(H)\} = 1 - \prod_{k \neq +} \frac{\lambda_+}{\lambda_+ - \lambda_k}. \quad (2)$$

The proof (Appendix A) holds for any multiplicities. Equation (2) specializes classical real and complex Gaussian quadratic-form theory [36], [37]. At the true AR(1) coefficient it is the CA law $1 - (1 + q^2/m)^{-m}$ [1], [3] applied to innovations. It evaluates coverage, the probability that the next sample falls inside the prediction disc, at any linear center and quadratic scale, texture and noise level, without simulation.

Corollary 1. With $\nu = 0$, (2) does not depend on the texture. With $\nu > 0$ texture invariance is generally lost [34], [38]; the CNR-conditional coverage is $c \mapsto G_{cR+I}(\hat{q})$, which tends to the pure-clutter value as $c \rightarrow \infty$ and to $G_I(\hat{q})$ as $c \rightarrow 0$.

Corollary 2. A Swerling-1 target $\sqrt{SP} g v$ (P the local clutter-plus-noise power, S the SCR relative to it, $g \sim \mathcal{CN}(0, 1)$, v the target signature over the episode) is detected with $P_d = 1 - G_{\Sigma+SP v v^\top}(\hat{q})$. For a target in Y only, pure AR(1) clutter and the true coefficient,

$$P_d = \left(1 + \frac{q^2}{m} \frac{1 - |\rho|^2}{1 - |\rho|^2 + S}\right)^{-m}, \quad (3)$$

the CA-CFAR law with the SCR multiplied by the whitening factor $1/(1 - |\rho|^2)$.

Because the clutter innovation powers are conditionally independent exponentials, the OS-CFAR law also holds at onset.

Corollary 3. With $\nu = 0$, AR(1) speckle and $r = \rho$, the OS score with the k -th smallest of m innovation powers has $\mathbb{P}\{\text{exceedance}\} = \prod_{i=0}^{k-1} (m-i)/(m-i+q^2)$ for every texture, the OS-CFAR law [2]; a target in Y only is detected with this probability, q^2 replaced by $q^2/\{1 + S/(1 - |\rho|^2)\}$.

B. Thermal noise caps the whitening gain

Proposition 2. Under (1) with AR(1) speckle, testing the prediction error against its own variance gains, over power-only detection and for a long history, $G_{\text{IN}}(c) = (1+c)/\{c(1+|r|^2-2\Re r\rho)+1+|r|^2\}$ for the IN-ARCP center and $G_{\text{opt}}(c) = (1+c)/V(c) \geq \max\{1, G_{\text{IN}}\}$ for the optimal predictor; with $V = \frac{1}{2}(\xi + \sqrt{\xi^2 - 4|\rho|^2})$ and $\xi = 1 + |\rho|^2 + c(1 - |\rho|^2)$.

The prediction error $Y - r H_m$ is a weighted two-pulse canceller, and G_{IN} is its clutter attenuation, which equals its improvement factor for a target in Y alone [4]. In noise-dominated cells, first-order normalization loses up to $10 \log_{10}(1 + |\rho|^2)$ dB, which the NA score avoids (proof in the Supplementary Material).

C. Persistent targets

A target that remains in the beam after its onset enters the history that normalizes it; in the limit of a strong target present in every sample, the score cannot detect it.

Proposition 3. For $m > 1$, $|r| < 1$ and a persistent target $A(e^{j\omega t})_{t=1}^{m+1}$ in every sample, the IN-ARCP score tends as $|A| \rightarrow \infty$ to

$$\begin{aligned} T_\infty &= \frac{\sqrt{m} |1 - r e^{-j\omega}|}{\sqrt{(1 - |r|^2) + (m-1)|1 - r e^{-j\omega}|^2}} \\ &\leq \sqrt{\frac{m}{m-1}}. \end{aligned}$$

Thus $P_d \rightarrow 0$ whenever $\hat{q} > \sqrt{m/(m-1)}$.

To follow a target that starts at a known pulse, we fix the model for the rest of this section: the speckle is AR(1) with coefficient ρ , $\nu = 0$ and $r = \rho$ (Theorem 1 needs none of




Fig. 2. Closed forms (lines) against Monte Carlo (markers; 10^5 episodes per point, 5×10^4 in (b)) at the true coefficient, AR(1) compound-Gaussian clutter with $\rho = 0.93e^{-0.2j}$, lognormal texture, $m = 16$, $P_{fa} = 0.01$. (a) Proposition 4; (b) horizon, Corollary 4; (c) OS scales (quadrature); (d) Theorem 1, P_{fa} against the hit rate for Gaussian interference 30 dB above the local power, Monte Carlo only (open downward triangles: no event in 10^5).

these assumptions). The innovations $\varepsilon_1 = \sqrt{1 - |\rho|^2} x_1$ and $\varepsilon_t = x_t - \rho x_{t-1}$ are then independent $\mathcal{CN}(0, \sigma_e^2)$ given the texture, and every statistic below is homogeneous of degree zero, so the texture cancels. A Swerling-1 target $\sqrt{SP} g e^{j\omega t}$, $g \sim \mathcal{CN}(0, 1)$ common to all pulses, that starts at pulse t_0 contributes $g\sqrt{S_w}$ to ε_{t_0} and $g\sqrt{S_w} d(\omega) e^{j\omega t}$ to ε_t , $t > t_0$, where

$$S_w = \frac{S}{1 - |\rho|^2}, \quad |d(\omega)|^2 = |1 - \rho e^{-j\omega}|^2, \quad (4)$$

$$S_w |d(\omega)|^2 = S w(\omega), \quad w(\omega) = \frac{P}{S_c(\omega)},$$

and $S_c(\omega) = P(1 - |\rho|^2)/|1 - \rho e^{-j\omega}|^2$ is the clutter spectrum. Under H_0 , whitening yields conditionally independent exponential innovation powers. It multiplies the SCR by $1/(1 - |\rho|^2)$ in the onset pulse and by $w(\omega)$ after it; for AR(p) speckle, $d(\omega)$ is the whitening filter's response. Self-masking is classical for clutter maps [14], [15]; to our knowledge, the post-onset law with its horizon, the OS law under persistence and the certificate are new in this form.

D. After the onset: a rank-one law and a visibility horizon

Proposition 4. Let the innovation-normalized score with m history innovations be thresholded at $q > 0$, and let a Swerling-1 target have whitened energy $\mathcal{E}_0$ in the tested innovation and $\mathcal{E}_H$ in the history innovations. With $\beta = q^2/m$, let $\lambda_{\pm}$ be the roots of

$$\lambda^2 - (1 + \mathcal{E}_0 - \beta - \beta \mathcal{E}_H) \lambda - \beta(1 + \mathcal{E}_0 + \mathcal{E}_H) = 0.$$

Then $\lambda_+ > 0 > \lambda_-$ and

$$P_d = \frac{\lambda_+}{\lambda_+ - \lambda_-} \left(1 + \frac{\beta}{\lambda_+}\right)^{-(m-1)}. \quad (5)$$

For a persistent target observed at look $1 \leq \ell \leq m - 1$, $\mathcal{E}_0 = S_w |d(\omega)|^2$ and $\mathcal{E}_H = S_w \{1 + (\ell - 1) |d(\omega)|^2\}$; at onset ($\ell = 0$), $\mathcal{E}_0 = S_w$, $\mathcal{E}_H = 0$, and (5) reduces to (3).

Corollary 4. For $1 \leq \ell \leq m - 1$, as $S \rightarrow \infty$ the look- ℓ detection probability tends to one if $\ell < \ell^*$ and to zero if $\ell > \ell^*$, where

$$\ell^*(\omega) = 1 + \frac{m}{q^2} - \frac{1}{|d(\omega)|^2}. \quad (6)$$

Since $\ell^* < 1 + m/q^2$, a strong target is kept for fewer than $1 + m/q^2$ post-onset looks, and not near clutter Doppler, where $|d(\omega)|^2 \approx (1 - |\rho|^2)$. At $m = 16$ and $P_{fa} = 0.01$, $m/q^2 = 3.0$; for $|\rho| = 0.93$ the horizon is 3.7 looks at opposite Doppler and negative at clutter Doppler (Fig. 2a,b). A guard delays scale contamination. Keep the tested energy $\mathcal{E}_0$ unchanged; set $\mathcal{E}_H = 0$ for $\ell \leq \Delta$ and use the history energy at $\ell - \Delta$ thereafter. While the history is partially contaminated, the critical value is $\Delta + \ell^*$; the false-alarm law is unchanged. Once $\ell \geq \Delta + m$, the entire history contains the target and $\mathcal{E}_H = S + S_w(m - 1) |d(\omega)|^2$, consistent with Proposition 3.

With the OS scale, the target occupies ℓ history innovations after the onset, and P_d is a two-dimensional integral (quadrature). For a strong target and $\ell \leq m - k$, the contaminated

powers lie above the scale, which becomes the k -th smallest of $m - \ell$ clean powers, and

$$P_d \simeq \prod_{i=0}^{k-1} \frac{m - \ell - i}{m - \ell - i + q^2/(1 + S_w|d|^2)}, \quad \ell \leq m - k, \quad (7)$$

while $P_d \rightarrow 0$ for $\ell > m - k$ provided $q^2 > \max\{1, |d(\omega)|^2\}$, as in the settings tested (Fig. 2c).

E. Dwell integration and self-normalized Doppler

Over a dwell, non-coherent integration of K innovation powers against the mean of m history powers, at threshold θ and with $\tau = \theta/m$, has

$$P_{fa} = \sum_{i=0}^{K-1} \binom{m+i-1}{i} \frac{\tau^i}{(1+\tau)^{m+i}}, \quad (8)$$

and a Swerling-1 target starting at the first dwell pulse, with clean history, whitened energy $\mathcal{E}_D = S_w\{1 + (K-1)|d|^2\}$, $\delta = 1 + \mathcal{E}_D$ and $\tau' = \tau(1 - 1/\delta)$, has

$$P_d = \sum_{i=0}^{K-2} \binom{m+i-1}{i} \frac{\tau^i}{(1+\tau)^{m+i}} + \left(1 - \frac{1}{\delta}\right)^{-(K-1)} \times \left[\left(1 + \frac{\tau}{\delta}\right)^{-m} - \sum_{i=0}^{K-2} \binom{m+i-1}{i} \frac{\tau'^i}{(1+\tau')^{m+i}}\right], \quad (9)$$

the CA non-coherent-integration law [3], [4] with the whitened energy. For clutter power summed over K correlated pulses, matching the conditional mean and variance gives $K_{\text{eff}} = K^2 / \sum_{i,j} |\rho|^{2|i-j|}$ effective looks (1.41 for $K = 8$ at $|\rho| = 0.93$).

The P-ANMF divides the periodogram of N whitened innovations by their energy. For a known on-grid Doppler and a target present before the dwell, the statistic is Beta(1, $N-1$) under H_0 for every texture, and

$$P_d = \left[1 + \frac{P_{fa}^{-1/(N-1)} - 1}{1 + NSw(\omega)}\right]^{-(N-1)}, \quad (10)$$

the normalized-matched-filter law [7], [19] with the whitened SCR. Maximizing over the N DFT bins gives Fisher's g statistic [39]; inclusion-exclusion gives its null law and, with one ordinate of increased mean, its Swerling-1 detection law (Supplementary Material). Dwell normalization avoids the history self-masking of Proposition 3.

F. Conformal calibration

For i.i.d. continuous scores and $k_\alpha \leq n$, the false-alarm probability given the n calibration scores is Beta($n+1 - k_\alpha, k_\alpha$) [18]. With ties, as for clipped or binary statistics, $\mathbb{P}\{T > \hat{q}\} \leq \alpha$ still holds, conservatively. Because k_α is a ceiling, $\mathbb{P}\{P_{fa} > 2\alpha\}$ is not monotone in n . It first falls below 0.05 at $n = 149, 1,497$ and $14,978$ for $\alpha = 10^{-2}, 10^{-3}, 10^{-4}$, and it stays below for every n beyond about $3.15/\alpha$. Our 10^{-4} calibration sets of 17,200–23,100 per-look episodes and 12,100–16,300 dwell segments give 0.02–0.09. These Beta calculations require independence; non-overlap alone does not

validate them for correlated radar episodes. For a known-scale Swerling-1 detector, the expected P_d under calibration is a ratio of Beta functions, from which the CFAR loss of a conformal threshold follows [40] (Supplementary Material).

G. Certified integration under pulsed interference

Pulsed interference (asynchronous pulses of another radar, jamming pulses or FMCW chirp crossings) corrupts a few pulses with interference of arbitrary, unknown power. We call such a corrupted pulse a *hit*, to distinguish it from an exceedance, a statistic above its threshold. For the whitened statistics, a hit at pulse t_h corrupts the innovations $t_h, \dots, t_h + p$. The clipped integrator $T = \sum_{i=1}^K \min\{|\varepsilon_{m+i}|^2/s_{k_H}^2, \kappa\}$ and its binary counterpart admit a certificate that does not depend on the interference power.

Theorem 1. *For a clean episode x and a set $\mathcal{J}$ of corrupted innovations, let $W(x, \mathcal{J})$ be T computed with the corrupted history powers set to zero and the corrupted dwell terms set to κ (binary: counted as exceedances). If this scale is zero, set $W = K\kappa$ (binary: K). Then $T \leq W(x, \mathcal{J})$ for every interference value, and W is nondecreasing in $\mathcal{J}$. Let $\hat{q}$ be the split-conformal quantile of $W(X_i, \mathcal{J}_i)$ over clean calibration episodes X_i , with masks drawn i.i.d. from $\mathcal{D}$, independently of all clean episodes and of the test hit set. If those episodes are exchangeable and the test hit set is independent of all clean episodes and dominated by $\mathcal{D}$ in the inclusion order (a coupling $\mathcal{J} \subseteq \tilde{\mathcal{J}} \sim \mathcal{D}$ exists), the false-alarm probability is at most α for any interference amplitudes, including adversarial ones.*

A bound on the hit rate alone does not give this dominance: periodic or bursty interference with the same rate can put more hits into one segment than independent hits do. A trimmed sum, which discards the largest dwell terms, uses the same bound with corrupted dwell terms set to $+\infty$. Its population threshold is infinite if the probability of more corrupted terms than discarded exceeds α ; finite-sample thresholds also depend on the sampled masks. We discard the fewest terms for which this probability is at most $\alpha/5$: 4 of eight at a 2% hit rate and $\alpha = 10^{-2}$, since a hit corrupts two AR(1) innovations. Fig. 2d checks the theorem in simulation.

V. EXPERIMENTAL DESIGN

IPIX sea clutter. Fourteen McMaster IPIX Dartmouth 1993 stare sessions [41]: X-band, 1 kHz pulse repetition, 14 range cells of 15 m, six dates with wave heights of 0.8–3.8 m [22]. Provider preprocessing is used; target cells and one guard cell on each side are excluded from clutter. Section VI-C uses the primary target cell. The two like-polarized channels are analyzed separately. Each session is split into thirds, separated by 2,000-pulse gaps, that serve as training, calibration and test data. One of the three assignments of these roles was used while the methods were developed; the other two give the 56 evaluation units (session $\times$ channel $\times$ assignment), which therefore reuse recordings seen during development and are not an untouched confirmation. Day transfer and two series held out until the end are in the Supplementary Material.

77 GHz FMCW data. The open JKU parking-lot dataset [42]: 76–77 GHz chirp sequences (34 μ s chirp period, 128 chirps per frame, 100 frames) recorded by two radar stations observing a static scene of parked cars, with a clean reference run and runs with an unsynchronized interferer, dense (scenario A, 45% of chirps hit) and sparse (scenario C, 4.6%). A further scenario-A run adds a walking person. A chirp counts as a hit when more than 10% of its range bins exceed ten times the reference run’s median power.

NetRAD S-band sea clutter. Fourteen recordings of the open NetRAD data [43], [44]: 2.4 GHz, 45 MHz chirp, 1 kHz pulse repetition, monostatic node, Cape Peninsula, 9 June 2011, seven HH and seven VV recordings of 130 s, with the providers’ clutter range cells (41–101 per recording). No statistic was computed on these data before their protocol was frozen. Detectors and settings are those used on IPIX, except that the ANMF secondary data come from the cells two to four cells away (15–30 vectors). The 28 units are recording $\times$ assignment, and intervals resample recordings. A pulse is flagged when more than 10% of its range bins exceed ten times the bin’s median power over the recording. Whether flagged pulses are interference or range-extended sea spikes is not established.

Camera-labelled pedestrians. Three UW 77 GHz records acquired in 2019 provide ADC frames and external camera labels [45], [46]. One fixed chirp/transmit/receive channel per frame supplies range-bin slow time; whole-record roles and camera-screened clutter masks are fixed before scoring (Supplementary Material).

Protocols. Methods and expected outcomes were recorded before computation. The Supplementary Material identifies data exposure, failed conditions and exploratory choices.

Detectors and thresholds. Per-look scores use $m = 16$. Dwell detectors use a 16-pulse history and an 8-pulse dwell that starts at the first full-amplitude target pulse (the same idealization for every detector), the median as OS history scale, clip level $\kappa = 6$ (4 to 18 in a clip-level comparison) and binary threshold $\eta = 6$. The ANMF uses secondary vectors from clutter cells at least two cells away, at five time positions (20–35 vectors). Power detectors normalize by the same 16 history samples (CA16) or by the cell’s power over ± 1024 pulses (CAloc, non-causal). Conformal thresholds are calibrated on clean episodes from the calibration third. For the low-false-alarm analysis these episodes do not overlap (17,200–23,100 per unit; 12,100–16,300 dwell segments). Model-based thresholds are the laws of Section IV at the fitted model, Gaussian plug-ins, the Kraut–Scharf and fixed-point ANMF laws, Gaussian simulation and an AR residual bootstrap of the sieve type [12].

Targets, interference and statistics. Swerling-1 targets are injected from -5 to 25 dB SCR relative to local power: at onset, persistent from a given pulse with random Doppler, or rising linearly over L pulses. Both IPIX test segments and secondary data receive independent pulse hits with probability $p_h \in \{0.02, 0.05\}$ and complex Gaussian amplitude 10 or 30 dB above local power. On the 77 GHz data, injected beat-tone targets pass through the same raw-sample processing. The 95% intervals use six-day cluster bootstrap resampling (10,000

TABLE I
MEASURED FALSE-ALARM RATE / DESIGN α ON IPIX TEST THIRDS (56 UNITS, POOLED). CONFORMAL: SPLIT-CONFORMAL THRESHOLD. MODEL-BASED: THE RULE IN PARENTHESES AT THE FITTED MODEL. BOLD: OUTSIDE $[0.5, 2]$. DAY-CLUSTER INTERVALS ARE IN SUPPLEMENTARY TABLE S2; AT 10^{-4} EACH UNIT CONTRIBUTES 1.2–2.2 EXPECTED FALSE ALARMS.

α	Conformal			Model-based rule		
	10^{-2}	10^{-3}	10^{-4}	10^{-2}	10^{-3}	10^{-4}
<i>Per look</i>						
IN-ARCP (CA law)	1.01	1.03	0.73	0.96	1.02	1.92
residual bootstrap				1.63	2.43	4.27
IN-AR(4) (CA law)	1.00	1.02	0.86	1.03	1.19	2.73
OS scale, $k = 8$ (OS law)	1.04	1.02	0.92	2.15	8.34	45.9
NA-AR(4) (Gaussian)	1.00	1.05	1.04	1.56	2.79	7.46
Power, CA (white law)	1.00	1.00	0.75	0.80	0.56	0.49
Power, OS (white law)	1.01	0.98	0.75	2.07	5.56	18.3
<i>Eight-pulse dwell</i>						
Integration (law (8))	1.02	1.05	1.02	2.02	4.35	14.3
PAMF-H, one Doppler (CA law)	0.99	1.00	0.89	1.97	3.66	9.88
P-ANMF, one bin (Beta)	1.10	1.14	1.10	1.78	3.05	5.47
P-ANMF, max (Fisher’s g)	1.04	1.04	1.47	4.36	13.1	49.6
ANMF-SCM (Kraut–Scharf)	1.06	1.04	0.85	2.02	3.57	8.07
ANMF-Tyler (FP law)	1.04	1.05	0.82	1.62	2.66	6.01
Clipped, OS scale (Gaussian sim.)	1.05	0.67	0.03	4.84	17.2	59.9

replicates). They are approximate with six clusters; leave-one-day-out ranges are in the Supplementary Material.

VI. RESULTS

A. False-alarm control across clutter power and down to 10^{-4}

In Table I, conformal thresholds keep the measured rate of twelve of the thirteen statistics within 0.98–1.14 times design at 10^{-3} and 0.73–1.47 times at 10^{-4} (IN-ARCP: 0.73 [0.54, 0.99]). The exception is the clipped integrator, which is conservative (0.67 and 0.03) because its statistic saturates at $K\kappa$ and ties at the threshold. At 10^{-2} some ratios lie slightly above one, such as 1.10 [1.01, 1.19] for the one-bin P-ANMF. The excess exceeds sampling noise and suggests non-exchangeability between calibration and test thirds [47]. Quantifying its coverage effect requires a bound on distribution drift [48].

Model-based thresholds, by contrast, fail in the tail. At 10^{-4} , 10 of the 11 model-based laws of the whitened statistics and the ANMF exceed twice the design rate, by up to 60-fold (bold entries in Table I). The ANMF laws fail already at 10^{-3} , at 3.57 (SCM) and 2.66 (Tyler) times design. The Kraut–Scharf law fails 2.8–3.4-fold at 10^{-2} even in simulated compound-Gaussian clutter (Supplementary Material). The SCM failure is therefore already present under the model, and need not come from any departure of real clutter from it. A residual bootstrap of real innovations still exceeds design 4.27-fold at 10^{-4} , consistent with dependence among an episode’s innovations, which independent resampling discards. Exploratory diagnostics distinguish power-dependent OS and Fisher’s- g failures from more uniform clipped and ANMF failures (Supplementary Material). The only law that stays within twofold is the CA law of the IN-ARCP score (1.02 at 10^{-3} ; 1.92 [1.26, 2.61] at 10^{-4} , still above design), plausibly because its scale averages all 16 history innovations.




Fig. 3. Detection of Swerling-1 targets injected into IPIX test episodes ($m = 16$, conformal thresholds, per-look $P_{fa} = 0.01$). (a) Target in the tested pulse (74,980 episodes in total). NA-AR(4): noise-aware score, 8.0 dB less SCR than CA16. Circles: law of Corollary 2, model fitted on each unit’s training third. Empirical and law curves are unit means; the text averages per-unit gains. (b) Persistent 10 dB target, random Doppler (56 units, 236 440 episodes), with guards $\Delta = 0, 8, 16$, through 48 looks. Solid: episode-weighted measured P_d ; dashed: episode-weighted law of Proposition 4 at fitted AR(1) coefficients and actual thresholds (exploratory); shading: 95% day-cluster intervals.

At design 0.01, conformal IN-ARCP rates range from 0.0067 to 0.0118 across texture quintiles, a max/min spread of 1.7 (2.7 at 10^{-4}). Without normalization the spread is 87.5; raw-history RMS normalization gives 9.8. Innovation normalization therefore gives approximate pivotality—a null free of clutter power—where clutter dominates noise. In the larger low-false-alarm calibration sets, the spreads are 1.31 for IN-ARCP and 1.06 for CA16; the latter is more pivotal but needs 7.0 dB more SCR.

On the 77 GHz data, thermal noise breaks the texture invariance as Corollary 1 predicts. From the lowest CNR class to the 10–20 dB class, IN-ARCP coverage at a design of 0.90 rises from 0.897 to 0.962, more than the law’s 0.932 (mean absolute deviation over all classes 0.010), and the false-alarm rate at 10^{-2} falls further than predicted (post hoc; Table IV).

B. Detection of new returns

Fig. 3a shows targets that appear in the tested pulse. At $P_d = 0.5$, IN-ARCP needs 7.0 dB [4.3, 9.0] less SCR than power detection on the same pulses (CA16) and 9.8 dB less than the long-window power detector (CALoc; 3.5–13.8 dB by day, post hoc). The law of Corollary 2 predicts P_d with mean absolute error 0.021 per unit/SCR. Thermal noise explains the shortfall from the clutter-only whitening factor (11.2 dB; Proposition 2). The Gaussian IN-ARCP threshold gives the same sensitivity at a similar false-alarm rate here, attributing the gain to whitening.

Over an eight-pulse dwell (SCR table in the Supplementary Material), PAMF-H is the most sensitive detector, and AR non-coherent integration is close to it. The ANMF (Tyler) needs at least 5.4 dB [1.9, 8.8] more SCR than PAMF-H; this is a lower bound, because PAMF-H reaches $P_d = 0.5$ at the lowest SCR tested. The P-ANMF pays for normalizing by a dwell that contains the target. The clipped integrator costs 2.0 dB [0.0, 2.9] against AR(1) integration at $\alpha = 0.01$, and

the binary integrator costs more. The power OS detector on the cell’s own history (clutter-map OS-CFAR) needs 10.6 dB [7.0, 12.7] more SCR than IN-ARCP at the first look, so the sensitivity comes from whitening, not from the order statistic. A Hann-windowed eight-pulse Doppler filter bank with range CA normalization needs 5.0 dB more SCR than integrated IN-ARCP. All whitening detectors lose sensitivity for targets at the clutter Doppler (PAMF-H by 4.3 dB).

The 77 GHz JKU walk provides radar-labelled onsets after fast-time interference suppression. At 200 ms look spacing, clutter is nearly uncorrelated ($|r| = 0.12$): IN-ARCP detects 0.30 within three frames versus 0.27 for the power clutter map. Only 7 event blocks occur, fewer than the eight required, so the comparison is descriptive. A further check uses 3 University of Washington camera-labelled walks [45], [46], at 33 ms spacing. Training $|r| = 0.02$ –0.04 fails the correlation criterion, and no primary clean-entry episodes occur. Mean presence association is 0.28 for guarded IN-ARCP versus 0.29 for guarded CA, with mixed background calibration. Camera labels establish presence rather than precise native-bin radar onset (Supplementary Material).

C. Persistent and emerging targets

A persistent target eventually enters its own reference history. Fig. 3b follows this process through 48 looks on 236 440 IPIX episodes; the extension also tests 1 075 080 NetRAD episodes (Supplementary Material). Increasing the guard delays the collapse predicted by Proposition 4, without preventing it. At 10 dB and random Doppler, the 16-pulse guard increases mean per-look P_d over looks 9–16 by 0.82 [0.76, 0.87] on IPIX and 0.92 [0.90, 0.94] on NetRAD relative to no guard. These paired cluster intervals support a longer useful interval, rather than indefinite persistent-target detection. Mean null-rate ratios across all looks and guards are 0.99–1.03 and 1.03–1.04, respectively.

TABLE II

REMEDIES FOR PULSED INTERFERENCE ON IPIX ($\alpha = 0.01$, 56 UNITS). INTERFERENCE IS INJECTED 30 dB ABOVE THE LOCAL POWER; EACH PULSE IS HIT WITH PROBABILITY p_h , AND BURSTS ARE EIGHT PULSES LONG AT $p_h = 5\%$. ENTRIES ARE THE MEASURED P_{fa}/α AND THE P_d OF A 10 dB PERSISTENT TARGET. TRIMMED: 4 (2%) AND 5 (5%) OF EIGHT TERMS DISCARDED. BOLD: $P_{fa}/\alpha > 1.2$.

Detector	None	$p_h = 2\%$		$p_h = 5\%$		Bursts
	P_d	P_{fa}/α	P_d	P_{fa}/α	P_d	P_{fa}/α
Clipped, clean threshold	0.86	1.10	0.85	1.28	0.85	3.98
Certified clipped, $\kappa = 6$	–	0.55	0.81	0.12	0.64	1.25
Certified clipped, $\kappa = 9$	–	0.53	0.81	0.14	0.71	1.66
Certified trimmed	–	0.57	0.79	0.23	0.72	3.60
Blanking, mask-matched	0.79	1.08	0.79	1.21	0.78	4.14

The earlier eight-look comparison found a guard onset cost of 0.9 dB for eight pulses and 1.2 dB for sixteen. At clutter Doppler the guard offers little help: whitening also attenuates the target. After full history contamination, Proposition 3 explains the low sensitivity for either guard. The median OS score delays contamination too, but costs 2.0 dB for abrupt targets, versus 0.69 dB predicted by Corollary 3; the guard is cheaper.

Targets that emerge gradually are harder for the per-look screen. When the target amplitude rises linearly over L pulses, reaching full amplitude at the first look, IN-ARCP's first-look gain over CALoc falls from 9.9 dB for an abrupt onset in the ramp comparison to 4.7 dB for $L = 8$. At $L = 16$ the target is barely detected ($P_d = 0.014$ at 10 dB), whereas PAMF-H over the dwell loses only 1.9 dB. The exact law reproduces the ramps (mean absolute error 0.016). The ramp fills the normalizing history, and with the eight-pulse guard the probability of detecting a 16-pulse ramp within eight looks at 10 dB rises from 0.03 to 0.87.

The IPIX reference target tests the opposite extreme. It is present throughout each recording, so it lies in every history and no guard helps. Proposition 3 explains it, and a pre-specified test measures it for IN-ARCP and the clipped integrator: both exceed on the target cell at 0.0004 at $\alpha = 10^{-3}$, below their clutter rates. Self-normalized whitened Doppler detects the target with $P_d = 0.22$ over 8 pulses and 0.40 over 256 pulses. Its measured clutter false-alarm rates are 1.1 and 1.5 times the design 10^{-3} . Its rate is also strongly texture-dependent (exploratory): the quintile spread is 4.4 at 10^{-2} and 23 at 10^{-4} , against 1.3 and 2.7 for IN-ARCP, consistent with the loss of texture invariance under additive noise (Corollary 1); the cause was not isolated. Range CA integration barely detects the target ($P_d \leq 0.007$), and the ANMF reaches 0.13–0.18 over 8–16 pulses. Different training and evaluation protocols make the long-dwell literature comparison descriptive (Supplementary Material). At 1,024 pulses, P-ANMF has a measured false-alarm rate of 2.6 times design; the i.i.d. Beta law does not guarantee control for these non-exchangeable windows.

D. Pulsed interference

With injected pulsed interference on IPIX at $\alpha = 10^{-2}$, non-coherent integration (AR(1) or AR(4)) and PAMF-H exceed

their design false-alarm probability 10- to 23-fold over two hit rates and two powers, because a dwell pulse corrupted by interference is an outlier that the history does not predict. Self-normalized detectors (P-ANMF, ANMF) stay within 1.21 times design but lose sensitivity.

Without interference at 10^{-3} , $\kappa = 6$ makes the clipped integrator conservative through saturation ($P_{fa}/\alpha = 0.67$). A pre-specified saturation rule picks $\kappa = 9$ in 17 of 56 units and is met by no tested level ($\kappa \leq 18$) in 16. Applied to all units (post hoc), $\kappa = 9$ gives $P_{fa}/\alpha = 0.90$ and raises P_d at 10 dB from 0.46 to 0.68, with negligible loss at 10^{-2} .

Table II compares the remedies with dense calibration (about 14,600 dwell segments per unit). The certified clipped integrator stays below design, consistent with Theorem 1 under its exchangeability and hit-model assumptions. The injected hits follow the calibration hit model, so this checks the implementation and the cost, not robustness. Its P_d at 10 dB falls from 0.86 with a clean threshold and no interference to 0.81 at a 2% hit rate and 0.64 at 5% (0.71 with $\kappa = 9$). The certified trimmed sum performs similarly, and so does certified binary integration (P_d within 0.03; exploratory, Supplementary Material). At 10^{-3} with 5% hits, the certificate is vacuous for every clip level: the certified threshold detects no target.

Pulse blanking, calibrated on clean episodes blanked at hit positions from the hit model, is cheaper but not certified. It keeps $P_d = 0.78$ at 5% hits and stays within 1.21 times design at 10^{-2} . At 10^{-3} with 5% hits it exceeds design 1.7- to 2.2-fold. Eight-pulse bursts expose the hit-model boundary (Table II): an independent-hit certificate loses false-alarm control because dominance fails; a burst-matched certificate controls false alarms but never detects. Blanking also exceeds design.

On the 77 GHz data the interference is real and occupies whole chirps. The hit model draws random 24-chirp blocks from the hit masks of frames 0–49; the test uses frames 50–99. In the sparse run the certified clipped integrator holds the false-alarm rate at 0.54 times design without fast-time mitigation, whereas non-coherent integration runs at 1.42 times design overall and 3.56 times in segments with one or two hit chirps. Against clipping with a clean threshold, the certificate costs 3.1 dB on hit-free segments and 6.4 dB overall. In the dense run, real interference mainly masks targets rather than raising false alarms (non-coherent integration 0.81 times design), and certification is nearly vacuous. Fast-time zeroing returns the uncertified detectors to their clean-reference false-alarm rates.

E. A second sea-clutter radar

Table III repeats the key tests on NetRAD, unused in development. At 1 kHz its clutter is more correlated ($|r| = 0.97$ –0.98, at or near the 0.98 clamp, versus an IPIX median 0.96). IN-ARCP gains ≥ 12 dB over raw-power OS at onset (Supplementary Material). The guard raises mean per-look P_d from 0.30 to 0.92 (law 0.92, exploratory) and within-eight-look detection of a 16-pulse ramp from 0.02 to 0.94.

Thresholds are refitted on NetRAD; Supplementary Table S3 gives recording-cluster intervals. All thirteen statistics are within 1.01–1.36 times design at 10^{-3} . At 10^{-4} , IN-ARCP, AR(4), OS and power scores stay within 1.38, but five

TABLE III

IPIX AND NetRAD S-BAND SEA CLUTTER, SAME CODE AND SETTINGS EXCEPT THE ANMF SECONDARY CELLS (SECTION V). ENTRIES: P_{fa}/α AT $\alpha = 0.01$ UNLESS STATED; GUARD ROW: PER-LOOK P_d AT 10 dB, MEAN OVER LOOKS 1–8; P_d IN PARENTHESES AT 10 dB. CONFORMAL ROWS: THE 12 STATISTICS OTHER THAN THE CLIPPED INTEGRATOR (CONSERVATIVE ON IPIX). INJECTED INTERFERENCE 30 dB ABOVE THE LOCAL POWER. FLAGGED PULSES: NetRAD HH, CERTIFICATE CALIBRATED ON THEIR MASKS.

	IPIX	NetRAD
Conformal, 10^{-3} (12 statistics)	0.98–1.14	1.01–1.36
Conformal, 10^{-4} (12 statistics)	0.73–1.47	1.02–4.26
Model-based laws $> 2\times$ at 10^{-4}	10 of 11	9 of 11
IN-ARCP CA law at 10^{-4}	1.92	1.70
Guard: P_d per look, $\Delta = 0 / 8$	0.24 / 0.82	0.30 / 0.92
5% hits: AR(1) integration	21	22
5% hits: certified clip (P_d)	0.12 (0.64)	0.08 (0.92)
5% bursts: blanking	4.14	4.11
Flagged pulses: integration / certified (P_d)	–	1.05 / 0.04 (0.04)

statistics exceed twice design: the noise-aware score (2.08), integration, PAMF-H and both P-ANMF statistics (up to 4.26). Exploratory attribution associates most integration and PAMF-H excess with two units whose calibration thirds lack flagged pulses (without them, 1.56 and 1.33). P-ANMF excess comes from one unit with flags in both thirds; fewer than half of its exceedances coincide with them. IN-ARCP's texture spread is 2.40 versus CA16's 1.51 and AR(4)'s 1.11, suggesting incomplete AR(1) whitening at the clamp.

The rule flags 3–23% of the HH pulses, in bursts of several pulses, yet uncertified non-coherent integration stays at 1.05 times design at 10^{-2} (1.20 at 10^{-3}). A certificate calibrated on these masks is nearly vacuous ($P_d = 0.04$), plausibly because the flagged pulses are frequent and come in bursts.

F. Acquisition with unknown crossing time

With unknown crossing time, a continuously migrating target crosses three native NetRAD bins during a 1,024-look search, calibrated at acquisition $\alpha = 0.01$. Across 24 fixed speed/SCR/response conditions (4256 acquisitions each), guarded IN-ARCP acquires only 0.004–0.030, with spatially local probability at most 0.019. P-ANMF has a higher point estimate in every condition; measured null-rate ratios are 1.62 and 0.94, respectively. The injection uses idealized sinc responses because the provider's complex range response is unavailable.

VII. DISCUSSION

A. Where the detector works and where it fails

Table IV links detector choice to correlation at the tested lag, noise, history contamination and pulse clustering. Sea-clutter calibration and ground-scene noise, interference and pedestrian tests address complementary conditions. False-alarm control and useful detection require separate evidence.

Exact innovation laws require true AR whitening, common texture and no thermal noise; the quadratic-form law allows arbitrary Gaussian covariance. Fitted laws predict measured detection curves closely but can fail in the far tail. Marginal conformal validity requires exchangeability; within-session

support is approximate, and online recalibration risks target contamination (Supplement). Certification also requires clutter-independent, dominated hit positions. Measured agreement and empirical hit masks do not prove these assumptions.

The step and ramp tests inject targets at known timing and dwell alignment. Real-target evidence remains descriptive: IPIX contains a continuously present target; the ground walks lack informative external onsets at strongly correlated lags. Calibration down to 10^{-4} remains unestablished beyond the two sea-clutter datasets.

Continuous migration can contaminate the scale before the target peak arrives; a guard does not ensure a clean history (Supplementary Material).

B. Design rules for a cell-wise detector

- 1) Whiten with an AR model fitted on clutter training data and normalize by the innovation scale of a window ending Δ pulses before the tested pulse ($\Delta = 8$ on IPIX and NetRAD; a longer guard extends the ideal horizon but its onset cost must be measured). Choose the look interval by measuring $|r|$ at that lag first: the gain follows the correlation, and vanishes between 77 GHz frames. Targets present throughout the history need self-normalized statistics (modest P_d on the real IPIX target).
- 2) Calibrate thresholds conformally on clean clutter: for i.i.d. continuous scores, beyond about $3.15/\alpha$ episodes, $\mathbb{P}\{P_{fa} > 2\alpha\} \leq 0.05$. At 10^{-4} that is 31,477 episodes, nominally 87–112 s of clean recording at 1 kHz across the 7–9 clutter cells of an IPIX file, ranked offline; dependence requires separate validation. Approximate pivotality permits a shared threshold where clutter dominates noise (texture spread 2.7 at 10^{-4}), at p complex multiply-accumulates per cell per pulse and 25 stored samples per cell.
- 3) Excise pulsed interference in fast time where possible. Otherwise, at low hit rates, integrate clipped innovations ($\kappa \approx 9$, post hoc) with a certified threshold and a hit model that dominates the hit pattern, not only its rate. The tested burst-matched certificate controls false alarms but detects no target; blanking exceeds design.

VIII. CONCLUSION

Under the true AR compound-Gaussian model, common texture and no thermal noise, innovation normalization recovers classical CFAR laws. The post-onset and OS laws quantify self-masking; 48-look injected trajectories confirm the finite benefit of guards. Refitted conformal thresholds improve IPIX/NetRAD tail calibration over the tested model laws, with NetRAD failures at 10^{-4} . JKU tests quantify noise-dependent calibration, real-interference control and its sensitivity cost. Fast-time excision restores clean-reference rates; weak frame correlation limits whitening gain in JKU/UW pedestrian checks. Certification requires exchangeability and dominated, clutter-independent hit positions. The tested burst certification and continuous-migration acquisition lack useful sensitivity. Operational validation still needs independently timed real-target onsets in correlated clutter and a measured range response.

TABLE IV

MEASURED OUTCOMES (P_{fa}/α , DEFAULT $\alpha = 10^{-2}$). IPIX UNLESS NAMED; INJECTED SCR 10 dB (MIGRATION: 0/10/20 dB). RAMP GUARDS: $\Delta = 0 \rightarrow 8$; PERSISTENT-TARGET GUARDS: $0 \rightarrow 16$. GAIN INTERVALS: 95% DAY-CLUSTER. EQUAL NOMINAL DESIGNS CAN HAVE DIFFERENT MEASURED FALSE-ALARM RATES. EVIDENCE: SUPPLEMENT TABLES S2–S3, S9, S11–S12, S14–S18, S20–S24; SECTIONS 13, 15, 17 AND 21.

Test	Measured outcome
IPIX, $\alpha = 10^{-4}$	Closer to design. Conformal: 12/13 at 0.73–1.47; clipped: 0.03. Model laws: 10/11 above 2.
NetRAD, refitted thresholds	Mixed tail accuracy. All 13 at ≤ 1.36 for 10^{-3} ; at 10^{-4} , IN-ARCP, AR(4), OS and power at ≤ 1.38 , 5 other statistics above 2.
Clutter power and CNR	Approximate power invariance. IPIX spread 1.7 versus raw-RMS 9.8. At 77 GHz, ratio 1.07 (< -5 dB CNR) to 0.08 (10–20 dB).
Abrupt onset, IPIX	Sensitivity gain. 7.0 dB [4.3, 9.0] less SCR than CA16 at $P_d = 0.5$; law's mean absolute P_d error 0.021.
Persistent target, guard	Longer visibility, eventual self-masking. Over looks 9–16, $\Delta = 16$ increases mean P_d by 0.82 on IPIX and 0.92 on NetRAD versus $\Delta = 0$. All guards subsequently lose sensitivity.
16-pulse ramp, guard	Improved detection within eight looks. IPIX 0.03 $\rightarrow$ 0.87; NetRAD 0.02 $\rightarrow$ 0.94.
Continuous migration, NetRAD	Guard insufficient. IN, $\Delta = 8$: acquisition 0.004–0.030 over all 24 fixed conditions; P-ANMF higher in each. Null ratios 1.62 and 0.94, respectively.
Clutter Doppler, IPIX	Low sensitivity. $P_d = 0.024$ at look 1 without a guard; guarded mean 0.07.
Persistent real IPIX target	IN-ARCP self-masks. P-ANMF: $P_d = 0.22$ –0.40 over 8–256 pulses; ratios 1.1–1.5 at nominal 10^{-3} .
UW camera-associated pedestrians	No demonstrated whitening advantage. Camera-labelled walks: mean presence association 0.28 (IN, $\Delta = 8$) versus 0.29 (CA); $ r = 0.02$ –0.04. No eligible primary clean entries.
Independent injected hits	Control at a detection cost. Uncertified integration: 10–23 times design. Certified clipping ($\kappa = 6$), IPIX: below design, $P_d = 0.81 / 0.64$ at 2% / 5% hits.
NetRAD, 5% injected hits	Matched-hit result reproduced. Certified ratio 0.08, $P_d = 0.92$; uncertified integration 22 times design.
5% hits, $\alpha = 10^{-3}$	No useful certified detection. Certificate: no target detected; blanking: up to 2.2 times design on IPIX.
Eight-pulse bursts, 5%	Hit-model boundary. Independent-hit certificate: 1.66 times design ($\kappa = 9$); burst-matched certificate: no detection; blanking: 4.1 times design on IPIX.
JKU real interference, 77 GHz	Fast-time excision preferred. Sparse certificate: ratio 0.54, 6.4 dB overall cost; dense: nearly vacuous. Zeroing restores clean-reference rates.
NetRAD flagged pulses	Little certified sensitivity. $P_d = 0.04$ for the certificate calibrated on frequent, clustered flags.

APPENDIX A

PROOF OF PROPOSITION 1

Write $X = \Sigma^{1/2}U$ with $U \sim \mathcal{CN}(0, I)$. The event is $\{X^H M_q X \leq 0\}$, and $X^H M_q X = \sum_k \lambda_k E_k$ with independent standard exponentials E_k . M_q is positive in the direction e_{m+1} , and $M_q + \frac{q}{m} \text{diag}(Q, 0)$ has rank one, so by Weyl's inequality M_q has exactly one positive eigenvalue; by Sylvester's law of inertia so does $\Sigma^{1/2} M_q \Sigma^{1/2}$. Unitary diagonalization preserves the independent standard complex-Gaussian coordinates of U . All other nonzero λ_k are negative. Conditioning on $Z = \sum_{k \neq +} |\lambda_k| E_k$ and using $\mathbb{P}\{E_+ > z\} = e^{-z}$,

$$\mathbb{P}\{\lambda_+ E_+ > Z\} = \mathbb{E} e^{-Z/\lambda_+} = \prod_{k \neq +} \frac{\lambda_+}{\lambda_+ - \lambda_k},$$

whatever the multiplicities, and (2) is its complement.

APPENDIX B

PROOFS OF COROLLARY 2, PROPOSITIONS 3 AND 4, AND THEOREM 1

Corollary 2. At $r = \rho$ in pure AR(1) clutter of power σ^2 , $Y - \rho H_m = \varepsilon + \sqrt{S} \sigma g$ with independent $\varepsilon \sim \mathcal{CN}(0, \sigma^2(1 - |\rho|^2))$, $g \sim \mathcal{CN}(0, 1)$, and $ms_\rho^2/(\sigma^2(1 - |\rho|^2))$ is a Gamma(m) variable; (3) is its moment-generating function.

Proposition 3. The score is homogeneous of degree zero, so its value at $C + Au$, with clutter C and $u = (e^{j\omega t})_t$, equals that at $C/A + u$ and tends to its value at u ; at u the numerator is $|1 - re^{-j\omega}|$ and $ms_r(u)^2 = (1 - |r|^2) + (m - 1)|1 - re^{-j\omega}|^2$. Dominated convergence gives $P_d \rightarrow 0$.

Proposition 4. In whitened units the episode is $\mathcal{CN}(0, I + vv^H)$, with $|v_{m+1}|^2 = \mathcal{E}_0$ and $\sum_{t \leq m} |v_t|^2 = \mathcal{E}_H$, and the

event is $x^H M x > 0$, $M = \text{diag}(-\beta I_m, 1)$. The eigenvalues of $M(I + vv^H)$ are $-\beta$ ($m - 1$ times) and the roots of $(\lambda - 1)(\lambda + \beta) - \mathcal{E}_0(\lambda + \beta) + \beta \mathcal{E}_H(\lambda - 1) = 0$ (rank-one update of a diagonal matrix). With $\lambda_+ > 0 > \lambda_-$, $\mathbb{P}\{\lambda_+ E_+ > |\lambda_-| E_- + \beta \Gamma_{m-1}\}$, with Γ_{m-1} a sum of $m - 1$ unit exponentials, is (5). For the horizon, $\lambda_+ \rightarrow \infty$ or stays bounded as $S \rightarrow \infty$ according to the sign of $|d|^2 - \beta\{1 + (\ell - 1)|d|^2\}$.

Theorem 1. The k_H -th smallest history power is nondecreasing in each power, so for any values of the corrupted entries it is at least its value with those entries set to zero. Uncorrupted dwell terms are then clean powers over a scale at least that large, and corrupted clipped terms are at most κ . Hence $T \leq W(x, \mathcal{J})$, and W is nondecreasing in $\mathcal{J}$. If $\mathcal{J} \subseteq \tilde{\mathcal{J}} \sim \mathcal{D}$ with $\tilde{\mathcal{J}}$ independent of the clutter, $T \leq W(X, \tilde{\mathcal{J}})$, and $(X_i, \mathcal{J}_i)$ and $(X, \tilde{\mathcal{J}})$ are exchangeable, so the split-conformal rank argument bounds the false-alarm probability by α . Details and the remaining laws are in the Supplementary Material.

CODE AND DATA AVAILABILITY

Code, protocols, results and figure scripts: <https://github.com/razaumair2203-ux/inarcip> (inarcip 0.3.0, tag r22-scope-submission). Data: McMaster IPIX [41], UCL NetRAD [43], JKU 77 GHz [42] and UW ADC [45].

ACKNOWLEDGMENT

The authors thank the McMaster IPIX group, M. Ritchie and the NetRAD team (UCL and University of Cape Town), L. Rienessl and R. Feger (JKU), and X. Gao and the UW team for open radar data. Anthropic Claude Code assisted with derivations, simulations, analysis code and text in all sections; OpenAI Codex assisted with code, mathematical/citation

checks and editing. The authors designed and validated the technique, experiments and test plan and take full responsibility for the content.

REFERENCES

- [1] H. M. Finn and R. S. Johnson, "Adaptive detection mode with threshold control as a function of spatially sampled clutter-level estimates," *RCA Review*, vol. 29, no. 3, pp. 414–464, 1968.
- [2] H. Rohling, "Radar CFAR thresholding in clutter and multiple target situations," *IEEE Trans. Aerosp. Electron. Syst.*, vol. AES-19, no. 4, pp. 608–621, 1983.
- [3] P. Gandhi and S. Kassam, "Analysis of CFAR processors in nonhomogeneous background," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 24, no. 4, pp. 427–445, 1988.
- [4] M. A. Richards, *Fundamentals of Radar Signal Processing*, 2nd ed. New York: McGraw-Hill Education, 2014.
- [5] F. S. Abbasi, U. I. Ahmed, and S. Ahmed, "Combating radar pulse jamming using clipping-based noncoherent pulse integration," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 54, no. 4, pp. 1783–1789, Aug. 2018.
- [6] K. Ward, R. Tough, and S. Watts, *Sea Clutter: Scattering, the K Distribution and Radar Performance*. Institution of Engineering and Technology, 2013.
- [7] E. Conte, M. Lops, and G. Ricci, "Asymptotically optimum radar detection in compound-Gaussian clutter," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 31, no. 2, pp. 617–625, 1995.
- [8] E. Ollila, D. E. Tyler, V. Koivunen, and H. V. Poor, "Complex elliptically symmetric distributions: Survey, new results and applications," *IEEE Trans. Signal Process.*, vol. 60, no. 11, pp. 5597–5625, 2012.
- [9] T. Bai, J. Tang, and Y. Xu, "Complex VAE with heavy-tailed likelihood for radar target detection in sea clutter," arXiv:2606.10540, 2026.
- [10] Y. A. Rouzoumka, J. Pinsolle, E. Terreaux, C. Morisseau, J.-P. Ovarlez, and C. Ren, "Out-of-distribution radar detection with complex VAEs: Theory, whitening, and ANMF fusion," arXiv:2601.18677, 2026.
- [11] Q. Qu, W. Liu, J. Wang, B. Li, N. Liu, and Y.-L. Wang, "Enhanced CNN-based small target detection in sea clutter with controllable false alarm," *IEEE Sensors J.*, vol. 23, no. 9, pp. 10 193–10 205, May 2023.
- [12] C. Qu, X. Wang, and J. Yin, "Sea clutter radar target detector based on autoregressive sieve bootstrap," *Digit. Signal Process.*, vol. 145, p. 104298, 2024.
- [13] C. Qu, L. Yang, L. Zhen, X. Wang, and J. Yin, "Robust adaptive CFAR detection using AR-sieve bootstrap and sample autocovariance," *Digit. Signal Process.*, vol. 168, p. 105574, Jan. 2026, art. no. 105574.
- [14] M. Lops and M. Orsini, "Scan-by-scan averaging CFAR," *IEE Proc. F*, vol. 136, no. 6, pp. 249–254, 1989.
- [15] M. Naldi, "False alarm control and self-masking avoidance by a biparametric clutter map in a mixed interference environment," *IEE Proc.-Radar Sonar Navig.*, vol. 146, no. 4, pp. 195–200, 1999.
- [16] J. Roman, M. Rangaswamy, D. Davis, Q. Zhang, B. Himed, and J. Michels, "Parametric adaptive matched filter for airborne radar applications," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 36, no. 2, pp. 677–692, 2000.
- [17] J. Lei, M. G'Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, "Distribution-free predictive inference for regression," *J. Amer. Statist. Assoc.*, vol. 113, no. 523, pp. 1094–1111, 2018.
- [18] V. Vovk, "Conditional validity of inductive conformal predictors," in *Proc. ACML*, vol. 25, 2012, pp. 475–490.
- [19] S. Kraut and L. Scharf, "The CFAR adaptive subspace detector is a scale-invariant GLRT," *IEEE Trans. Signal Process.*, vol. 47, no. 9, pp. 2538–2541, 1999.
- [20] F. Pascal, Y. Chitour, J.-P. Ovarlez, P. Forster, and P. Larzabal, "Covariance structure maximum-likelihood estimates in compound Gaussian noise: Existence and algorithm analysis," *IEEE Trans. Signal Process.*, vol. 56, no. 1, pp. 34–48, Jan. 2008.
- [21] J. Michels, B. Himed, and M. Rangaswamy, "Evaluation of the normalized parametric adaptive matched filter STAP test in airborne radar clutter," in *Record of the IEEE 2000 International Radar Conference*. IEEE, 2000, pp. 769–774.
- [22] S. Haykin and X. B. Li, "Detection of signals in chaos," *Proc. IEEE*, vol. 83, no. 1, pp. 95–122, 1995.
- [23] J. Xu, J. Yu, Y.-N. Peng, and X.-G. Xia, "Radon-fourier transform for radar target detection, I: Generalized doppler filter bank," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 47, no. 2, pp. 1186–1202, 2011.
- [24] S. Bates, E. Candès, L. Lei, Y. Romano, and M. Sesia, "Testing for outliers with conformal p-values," *Ann. Statist.*, vol. 51, no. 1, 2023.
- [25] T. Bultingaire, I. Meraoumia, C. Kervazo, L. Denis, and F. Tupin, "Unsupervised radiometric change detection from synthetic aperture radar images," in *Proc. 32nd European Signal Processing Conference (EUSIPCO)*, 2024, pp. 2217–2221.
- [26] X. Meng, "Modified rank-sum nonparametric CFAR in multiple-target situation," *Journal on Advances in Signal Processing*, vol. 2026, 2026, art. no. 65.
- [27] S. D. Himonas, "CFAR integration processors in randomly arriving impulse interference," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 30, no. 3, pp. 809–817, Jul. 1994.
- [28] M. Toth, P. Meissner, A. Melzer, and K. Witrisal, "Performance comparison of mutual automotive radar interference mitigation algorithms," in *Proc. 2019 IEEE Radar Conference (RadarConf)*. IEEE, Apr. 2019, pp. 1–6.
- [29] M. Rameez, M. Dahl, and M. I. Pettersson, "Autoregressive model-based signal reconstruction for automotive radar interference mitigation," *IEEE Sensors J.*, vol. 21, no. 5, pp. 6575–6586, Mar. 2021.
- [30] A. Gendler, T.-W. Weng, L. Daniel, and Y. Romano, "Adversarially robust conformal prediction," in *International Conference on Learning Representations (ICLR)*, 2022.
- [31] S. H. Zargarbashi, M. S. Akhondzadeh, and A. Bojchevski, "Robust yet efficient conformal prediction sets," in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, vol. 235. PMLR, 2024, pp. 17 123–17 147.
- [32] A. Levine and S. Feizi, "Robustness certificates for sparse adversarial attacks by randomized ablation," *Proc. AAAI Conf. Artif. Intell.*, vol. 34, no. 4, pp. 4585–4593, 2020.
- [33] M. Zaffran, A. Dieuleveut, J. Josse, and Y. Romano, "Conformal prediction with missing values," in *Proc. 40th International Conference on Machine Learning (ICML)*, vol. 202, 2023, pp. 40 578–40 604, arXiv:2306.02732.
- [34] S. Watts, "Radar detection prediction in K-distributed sea clutter and thermal noise," *IEEE Trans. Aerosp. Electron. Syst.*, vol. AES-23, no. 1, pp. 40–45, 1987.
- [35] N. Ignatiadis and B. Sen, "Empirical partially Bayes multiple testing and compound χ^2 decisions," *Ann. Statist.*, vol. 53, no. 1, 2025.
- [36] J.-P. Imhof, "Computing the distribution of quadratic forms in normal variables," *Biometrika*, vol. 48, no. 3–4, pp. 419–426, 1961.
- [37] T. Y. Al-Naffouri, M. Moinuddin, N. Ajeeb, B. Hassibi, and A. L. Moustakas, "On the distribution of indefinite quadratic forms in Gaussian random variables," *IEEE Trans. Commun.*, vol. 64, no. 1, pp. 153–165, 2016.
- [38] F. Gini, "Sub-optimum coherent radar detection in a mixture of K-distributed and Gaussian clutter," *IEE Proc.-Radar Sonar Navig.*, vol. 144, no. 1, pp. 39–48, 1997.
- [39] R. A. Fisher, "Tests of significance in harmonic analysis," *Proc. R. Soc. Lond. A*, vol. 125, no. 796, pp. 54–59, Aug. 1929.
- [40] S. Watts, K. Ward, and R. Tough, "CFAR loss and gain in K-distributed sea-clutter and thermal noise," in *Proc. IET International Conference on Radar Systems*, 2007.
- [41] McMaster University IPIX Radar Group, "The IPIX radar database: Dartmouth 1993," <https://soma.ece.mcmaster.ca/ipix/dartmouth/data/sets.html>, accessed 29 Sep. 2026.
- [42] L. Rienessl and R. Feger, "77-GHz FMCW radar network ADC data under various interference scenarios at parking lot," Zenodo dataset, doi:10.5281/zenodo.21933926, 2026.
- [43] M. Ritchie, "NetRAD – monostatic & bistatic sea clutter dataset," University College London, dataset, doi:10.5522/04/32676582, 2026.
- [44] F. Fioranelli, M. Ritchie, H. Griffiths, S. Sandenbergh, and M. Inggs, "Analysis of polarimetric bistatic sea clutter using the NetRAD radar system," *IET Radar Sonar Navig.*, vol. 10, pp. 1356–1366, 2016.
- [45] X. Gao, G. Xing, Y. Luo, S. Roy, and H. Liu, "Raw ADC data of 77GHz MMWave radar for automotive object detection," 2022, IEEE DataPort dataset, doi:10.21227/xm40-jx59.
- [46] X. Gao, G. Xing, S. Roy, and H. Liu, "RAMP-CNN: A novel neural network for enhanced automotive radar object recognition," *IEEE Sensors J.*, vol. 21, no. 4, pp. 5119–5132, 2021.
- [47] M. Greco, P. Stinco, F. Gini, and M. Rangaswamy, "Impact of sea clutter nonstationarity on disturbance covariance matrix estimation and CFAR detector performance," *IEEE Trans. Aerosp. Electron. Syst.*, vol. 46, no. 3, pp. 1502–1513, 2010.
- [48] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, "Conformal prediction beyond exchangeability," *Ann. Statist.*, vol. 51, no. 2, 2023.